

Cavefish: Communication-Optimal Light Client Protocol for UTxO Ledgers

Anonymous author

Anonymous affiliation

Abstract

Blockchain light clients (LCs) are agents with limited computational or storage resources who cannot maintain a fully validated local copy of the ledger state. Instead, they rely on service providers (SPs), typically full nodes, to access data required for tasks such as constructing transactions or interacting with off-chain applications.

In this work, we introduce Cavefish, a novel protocol for UTxO-based platforms that enables LCs to interact with the ledger and submit transactions with minimal trust, storage, and computation. Cavefish defines a two-party computation protocol between an LC and an SP, in which the LC specifies a transaction and the SP constructs it. Consequently, the LC only receives a blinded version of the transaction, preventing it from modifying or reusing the transaction while still being able to verify that the transaction matches the original intent of the LC. The SP is compensated inside the constructed transaction, eliminating the need for another protocol or exchange.

To support this, we propose a variant of the predicate blind signature (PBS) scheme of Fuchsbauer and Wolf (Eurocrypt 2024), allowing the SP to obtain a valid signature on the unblinded transaction, which it can then broadcast on the network and post on chain. Moreover, the resulting signatures verify as standard Schnorr signatures. Our construction achieves a trustless interaction in which the LC achieves their transaction goal, and the SP receives fair compensation for their effort. When Cavefish is combined with hierarchical deterministic (HD) wallets, the LC can provide a single public key and chain code to the SP, reducing communication footprint to a minimum.

To further optimize communication and computational overhead, our PBS variant relaxes the unlinkability guarantees of traditional blind signatures in favor of efficiency. We argue that this relaxation is adequate, since transactions only need to be kept private until posted on a public ledger. We implement and benchmark the Non-interactive Argument of Knowledge (NArg) component of our protocol on two major UTxO-based blockchains. Despite being the most computationally demanding part, our results show that proving and verification times, as well as circuit sizes, are practical for real-world deployment.

2012 ACM Subject Classification Security and privacy → Security protocols

Keywords and phrases light clients, protocols, wallets, blockchain

Digital Object Identifier 10.4230/LIPIcs.CVIT.2016.23

1 Introduction

Blockchain technologies have emerged as a foundational component of decentralized systems, offering strong guarantees of data integrity, censorship resistance, and fault tolerance through cryptographic protocols and distributed consensus. Within this domain, the Unspent Transaction Output (UTxO) model represents a distinctive paradigm for managing asset ownership and validating transactions. The UTxO model was initially introduced by Bitcoin and subsequently adopted by other platforms such as Cardano. In contrast to account-based models, UTxO-based blockchains accommodate parallelism and concurrent processing more effectively but also introduce challenges in terms of complexity and client verification.

Full nodes in a UTxO-based blockchain are required to download and validate the entire chain history to ensure correctness and security. This requirement presents a significant barrier to participation for resource-constrained devices, such as smartphones and embedded systems. Light client (LC) protocols aim to mitigate this issue by enabling nodes to interact with the blockchain in a secure and efficient manner without maintaining full historical data. These

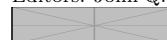


© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:30



Leibniz International Proceedings in Informatics



Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

47 protocols must strike a careful balance between minimizing resource consumption and pre-
 48 serving critical security properties, such as transaction inclusion, double-spending resistance,
 49 and above all, chain validity.

50 Blockchains are append-only data structures that grow continuously over time. As the
 51 chain length increases, it becomes prohibitively expensive for a light client (LC) to scan the
 52 entire history to verify past transactions or to locate a specific UTxO.

53 The question answered by this paper is “how can a user of a light client engage with a full
 54 node acting as a service provider to request and approve transactions (e.g. from their wallet) in
 55 a secure way without knowing anything about the current chain and ledger state and minimal
 56 communication effort?” In this paper, we present a solution to this question in the form of
 57 *Cavefish*, a novel intent-based light client protocol designed for UTxO-based blockchains.
 58 Cavefish enables LCs to avoid querying the current ledger state and submit transactions
 59 requiring only minimal local storage and computation. In order to submit a transaction on
 60 chain, the LC engages with a service provider (SP) in a two-party computation protocol that
 61 yields a signed transaction indistinguishable from one created by a full node. In addition to
 62 the low storage and computational requirements, our LC protocol is communication-optimal.
 63 After the LC has instructed the SP about the type of transaction it wishes to create, the
 64 protocol can be completed in as few as two rounds. The LC does not need to download let
 65 alone parse the blockchain. The only information the SP must obtain from the LC are the
 66 addresses where the funds are located. Cavefish is compatible with hierarchical (HD) wallets
 67 [49] allowing straightforward address discovery over a range of child addresses with the LC
 68 sending only a single public key together with its chain code. The only complexity is having
 69 the LC sign the requested signature after the SP constructs it. Sending the signature in the
 70 open would allow the LC to modify the transaction, potentially removing SPs fee. Instead,
 71 the SP transmits the result in a redacted form, which we call an *abstract transaction*.

72 The LC and SP then complete a blind signature protocol where the LC verifies if the
 73 transaction satisfies its defined specifications and only then creates valid signature(s) which are
 74 sent back to the SP. The abstract transaction is used to speed up this check. As a last step, the SP
 75 attaches the signatures to the transaction and posts them on the blockchain on behalf of the LC.

76 To make our scheme viable in the real world, we adapt the blind predicate Schnorr signature
 77 scheme from [24]. More precisely, we do not require the unlinkability requirement once the
 78 transaction has been published on the blockchain. Accounting for timing, values and payees
 79 the potential anonymity set for a blinded transaction is trivially small, and in any case the
 80 transaction signed by the LC only needs to stay *private until posted*. This insight allows us
 81 to introduce the notion of a *weakly* blind predicate signature scheme, a simplification that
 82 reduces the space and time complexity of the zero-knowledge component which asserts that
 83 the abstract transaction meets the LC’s specifications.

84 In addition to the low communication overhead, our light client protocol gives the SP the
 85 ability to be reimbursed for its computation time required to construct the transaction without
 86 the need of an additional protocol or exchange. The requested transaction includes the SP’s
 87 fee, i.e., an additional UTxO output, which makes up for the SP’s costs and small reward.
 88 Additionally, if the client wishes to engage with the SP over a period of time, our mechanism
 89 can be used to initialize a payment channel that can be used subsequently for fast payments
 90 to the SP without impacting the size of transactions beyond the first.

91 To summarize, our contributions are:

- 92 ■ We introduce the notion of intents, describing the desired end result of the light client’s
 93 actions, as opposed to the method by which they are accomplished. We believe this to
 94 be an important abstraction as it assists in the conception, development and study of

95 specialized efficient protocols as opposed to aiming for parity with full clients. To this end,
96 we describe a domain-specific language (DSL) allowing the concise construction of intents
97 for UTxO-based ledgers.

- 98 ■ We propose Cavefish, a light client protocol that avoids any enrollment or synchronizing
99 with chain whilst maintaining safety, providing compensation to the SP, imposing minimal
100 communication overhead and being compatible with any UTXO blockchains using Schnorr
101 signatures.
- 102 ■ We introduce the notion of weakly blind predicate signatures, motivated by our “private
103 until posted” goal. This brings together the notions of blind predicate signatures of [24]
104 with the notion of signatures on committed messages (SBCM) from [7].
- 105 ■ We implement and benchmark the zero-knowledge proof component of Cavefish for two
106 major blockchain platforms, Bitcoin and Cardano, and show that an “unoptimized” imple-
107 mentation achieves proving and verification times that are viable in a real-world deployment.

108 2 Background and Related Work

109 2.1 Light clients, off-chain payments, and chain explorer services

110 **Light clients.** Our approach facilitates transaction submission for a light client that is not
111 aware of the history of a blockchain. Instead, it operates more like a single-chain intent-resolution
112 protocol. Therefore, we compare our work with existing solutions for light clients as well as mech-
113 anisms that allow the interaction with a blockchain when in resource-constrained operation. Our
114 approach is different from a more “traditional” (i.e., what commonly is referred to) light client
115 that has the primary goal of syncing to the blockchain in order to acquire the information neces-
116 sary to interact with a smart contract or to submit a transaction. We first describe the common
117 idea of a light client and then outline concepts related to our work that complement light clients.

118 There is existing work analyzing light client functionality [2] [15]. The majority of light client
119 designs include the following main functionalities it is expected to perform: (1) issue queries, such
120 as for the balance of an account, or the state of a transaction, and (2) safeguard secret informa-
121 tion and submit transactions to the blockchain. In order to implement these functionalities, light
122 clients use several generic techniques, most notably: header verification and consensus evolution
123 verification. Unlike a *full node*, a header verification light client *only verifies the headers of blocks*,
124 and skips the verification of transactions and account balances [2]. The light client is able to trust
125 the resulting chain state if the data is signed by a sufficiently large set of full nodes (a multi-prover
126 approach). The technique was made popular with SPV in Bitcoin [37] and nearly universally
127 adopted by most practical approaches for light clients, e.g., [42], along with its variations [41]. Be-
128 cause the validator set can change, consensus evolution verification is needed for blockchains run-
129 ning on proof-of-stake. Another common technique is to compress the blockchain and/or ledger
130 state in order to reduce the information a light client need for functioning reliably (see, e.g., [9]).

131 Some light clients use game-theoretic approaches to eliminate the need for a trusted com-
132 mittee (a single-prover model), e.g. by implementing the slashing of previously deposited
133 collateral in case of misbehavior [34]. Our model is both single-prover and does not require
134 a deposit to be made, as neither the service provider nor the light client risk their assets during
135 the protocol execution. Moreover, our design is independent of the underlying consensus,
136 relying, instead, on the ledger model.

137 The main cryptographic building blocks that are used to realize those techniques are succinct
138 representation and proofs, such as data accumulators (often Merkle trees) and commitments
139 and SNARKs (e.g. used to aggregate multi-signatures [46]). Suitable signatures and hash

¹⁴⁰ functions are needed for certain light client designs. Examples of these include aggregate
¹⁴¹ signatures and threshold signatures [11, 2].

¹⁴² **Intents.** We also consider *intents* and *solver networks* to be related work. These designs
¹⁴³ attempt to establish a relationship between solvers and users via their (light) clients [20]. A
¹⁴⁴ light client issues an intent via as an abstracted transaction object. Then, solvers process the
¹⁴⁵ intent, incentivized by transaction fee or intent execution reward. The users are then free
¹⁴⁶ to accept or reject a solver’s proposal. In case the solvers are required to provide a deposit,
¹⁴⁷ slashing incentivizes honest behavior of a rational actor. Since concepts around solver networks
¹⁴⁸ are relatively new, there is currently no universal standard governing the specification for
¹⁴⁹ intents and abstract transaction objects. To the best of our knowledge, this is the first work
¹⁵⁰ describing an intent-based light client protocol taking advantage of specific features of the
¹⁵¹ UTxO ledger model that both minimizes communication and offers built-in incentive structure.

¹⁵² **Cross-chain intents.** A lot of work on intents applies is focused on building cross-chain in-
¹⁵³ tent resolution, which requires more sophisticated protocols that are able to communicate with
¹⁵⁴ multiple chains at once [50] . Some protocol designs choose to resolve intents on-chain directly
¹⁵⁵ [43], while other offer both layer-1 and layer-2 options [17]. Yet another designs is a consensus pro-
¹⁵⁶ tocol relying on minimal trust assumptions while conforming to the interledger standard [18].
¹⁵⁷ Focusing on a single chain allows us to construct a protocol which makes use of the unique features
¹⁵⁸ of the underlying blockchain, and requires reduced communication (only between two parties).
¹⁵⁹ In its simplicity and limited communication, our protocol also resembles payment channels.

¹⁶⁰ **Payment channels.** The concept of *payment channels*, and the related notion of *state*
¹⁶¹ *channels*, also bear similarities to our approach. Payment channels are a type of off-chain
¹⁶² mechanism for blockchains. A payment channel can be used to conduct a series of transactions
¹⁶³ without interacting with the main blockchain, e.g. [32] [44] [42]. The creation of a payment
¹⁶⁴ channel between parties requires locking funds in an on-chain transaction [35]. Some channels,
¹⁶⁵ such as Hydra [14], not only allow simple payments between users, but can simulate the majority
¹⁶⁶ of the on-chain transaction processing mechanism internally. This type of channel is usually
¹⁶⁷ called a *state channel*.

¹⁶⁸ Cavefish interacts well with payment channels: on the one hand cavefish can be used to
¹⁶⁹ setup a payment channel from a light client, and on the other setting up a payment channel
¹⁷⁰ allows the SP to be compensated offchain eliminating the on-chain overhead of Cavefish.

¹⁷¹ **API and Explorer Services.** Many realistic UTxO ledger implementations (e.g. Ergo
¹⁷² ¹, Cardano ², and BitCoin ³) are set up in a way that an invalid transaction will, in most cases,
¹⁷³ not result in any update to the ledger state, but get rejected instead. This makes services such
¹⁷⁴ as blockchain explorers ⁴ or Blockfrost (API as a service for accessing the Cardano blockchain)
¹⁷⁵ ⁵ some of the strongest competitors with our proposal, as they provide the data needed for
¹⁷⁶ the LC to construct their transaction, often with a relatively high degree of reliability. Unlike
¹⁷⁷ our design, these services do not currently appear to support transaction construction for light
¹⁷⁸ clients. Revenue from explorer services is either ad-based, or they may charge users in fiat
¹⁷⁹ currency, or the service is free in order to promote use of the specific service/blockchain. We,
¹⁸⁰ on the other hand, propose a protocol in which service providers are compensated in assets
¹⁸¹ on the same blockchain as the one to which their transaction gets applied, and the payment
¹⁸² structure consists of a one-time, no-setup atomic exchange of assets for services.

¹ <https://ergoplatform.org/>

² <https://cardano.org/>

³ <https://bitcoin.org/>

⁴ <https://beta.explorer.cardano.org/>

⁵ <https://blockfrost.dev/>

183 2.2 Hierarchical deterministic wallets and address discovery

184 UTXO based blockchains, including Bitcoin and Cardano use the concept of Hierarchical
 185 Deterministic (HD) wallets [49] where child addresses can be created in a deterministic way using
 186 a public key of a keypair sk, pk , and some short auxiliary data (the chain code). Given a public key
 187 pk , chain code c and requested index i for a child key, a hash function is used to produce a scalar
 188 s_i enabling the derivation of the child key as $sk_i = sk + s_i$. It is also possible to use s_i without
 189 knowing sk to directly derive the corresponding public key $pk_i = pk \cdot g^{s_i}$. The same method
 190 is used to produce a child chain code c_i , so that multiple generations of children are possible.

191 In our application, we can assume that a wallet can give the service provider (SP) a single
 192 public key and chain code, letting the SP do address derivation and lookup on their side using
 193 some agreed-upon bounds or heuristics on the index depth. This keeps the communication
 194 cost between the light client and the SP constant.

195 2.3 Schnorr blind signatures

196 As soon as their patent expired in 2008, Schnorr signatures [39] have been gaining significant
 197 adoption, replacing RSA in many scenarios due to their smaller size and faster verification.
 198 Compared to (EC)DSA, Schnorr offers similar efficiency and guarantees, but (EC)DSA se-
 199 curity proofs are most likely not possible without strong idealization [29]. As a consequence,
 200 EdDSA [6], a popular variant of the Schnorr scheme, is currently under consideration by NIST
 201 for standardization. The security of Schnorr signaatures is based on the discrete logarithm
 202 assumption [38], with proofs in the random oracle model (ROM) as well as in the algebraic
 203 group model (AGM) and the ROM [23].

204 Standard Schnorr signatures are now widely used in blockchains such as Bitcoin, Bitcoin
 205 Cash, Litecoin, and Polkadot. Monero, Zcash, and Cardano use the EdDSA variant. The
 206 adption of Schnorr and EdDSA signatures is driven by privacy [8] and scalability gains [36]
 207 but also the straightforward extension to *blind Schnorr signatures* [16]. Most Schnorr-based
 208 blind signature schemes require multiple rounds [16], which can make the protocol susceptible
 209 to denial-of-service attacks. Research has thus been focusing on making blind signatures
 210 *concurrently* secure where more than one session can be intertwined, such as [4, 12, 45, 28].

211 Our construction hinges on the the concurrently secure blind and partially blind signing pro-
 212 tocol for standard Schnorr signatures found in [24] by Fuchsbauer, et al.—the first work stating
 213 rigorous security guarantees for a practical blind signature scheme based on Schnorr signatures.
 214 Unlike schemes that had been presented before, [24] is not vulnerable to the attacks described
 215 in [47, 5] which showed that the hardness-assumption⁶ existing schemes relied on for concurrent
 216 security can be solved in polynomial-time. As our light client protocol is based on [24], it also
 217 provides partial and predicate blindness, two properties that allow us to describe a transaction
 218 in an abstract way featuring “redacted” parts. We briefly outline the construction of Fuchsbauer
 219 et al. [24] in the following. The scheme is equivalent to the blind signature scheme [16] by Chaum
 220 and Pedersen but, in order to make the protocol concurrently secure, adds a commitment phase
 221 at the beginning where the user sends an encrypted version of the message m and blinding
 222 values (α, β) to the signer. In addition to that, the second message by the user includes a
 223 zero-knowledge proof alongside the Schnorr challenge c . The zero-knowledge proof asserts to the
 signer that the initial (encrypted) commitment and c have been derived from the same m, α, β .

⁶ The ROS (Random inhomogeneities in a Overdetermined Solvable system of linear equations) problem was initially studied by Schnorr in [40]

23:6 Cavefish

225 In addition to obtaining concurrent security, one can leverage the zero-knowledge proof to
226 assert additional facts, most notably, the user can prove to the signer that certain predicate(s)
227 over message m holds. This effectively turns the fully blind scheme into a predicate blind
228 scheme. Furthermore, support for predicate blindness implies partial blindness [24] since
229 constructing a predicate that checks equality for parts of m can be used to assert to the user
230 that parts of the message correspond to the expected value(s).

231 3 Technical Background

232 **Notation.** We use $y \leftarrow x$ to denote that variable y is assigned the (possibly randomized)
233 evaluation of x . When x is a set, we denote uniformly random sampling from the set. We use
234 $a = b$ to denote boolean comparison between a and b and $z := w$ to denote equality by definition.
235 We use \parallel to denote concatenation of bitstrings. When algorithms are randomized, we denote
236 them as $A(x;r)$, where x is the input and r is the randomness, belonging to a randomness space
237 \mathcal{R} . When we write $y \leftarrow A(x)$ we imply $r \leftarrow \mathcal{R}; y \leftarrow A(x;r)$. In algorithm descriptions we write
238 $x,y,z \subseteq a$ to imply parsing a to obtain x,y,z .

239 Our protocols operate in the discrete log setting, where we assume the discrete logarithm
240 problem is hard. We also follow standard conventions with regards to public key encryption
241 using $(\text{PKE.KeyGen}, \text{PKE.Enc}, \text{PKE.Dec})$ to describe schemes and require IND-CPA security.
242 For signatures, we use Schnorr signatures $(\text{SDS.Setup}, \text{SDS.KeyGen}, \text{SDS.Sign}, \text{SDS.Ver})$. Finally,
243 we use parametrized non interactive Arguments $(\text{NArg.Rel}, \text{NArg.Setup}, \text{NArg.Prove}, \text{NArg.Ver},$
244 $\text{NArg.SimProve})$ to allow the relation being proven to depend on the group setting. For reasons
245 of space we present the full description of these primitives in Appendix A

246 The security of Schnorr signatures, has been well studied in the random oracle model
247 (ROM), where a hash replaced with an ideal random function the adversary can only call as
248 an oracle. However in this work the need arises to instantiate the underlying hash function H
249 so that we may reason about it within a proof system [3] e.g. prove that a known y is $y = H(x)$
250 for some secret x . This renders ROM-based proofs inapplicable necessitating an assumption
251 on the security of the scheme. This does not differ significantly from typical implementation
252 practices (where the hash function is infact drawn from a small pool for standardized options),
253 and also from the treatment of this issue in the literature [24].

254 ▶ **Definition 1.** A hash function generator $HGen(n) \rightarrow H$, on input $n \in \mathbb{N}$ generates a hash
255 function $H: \{0,1\}^* \rightarrow \mathbb{Z}_n$.

256 ▶ **Assumption 1.** [24] There exists a group generator $GGen$ and hash function generator $HGen$
257 s.t. the Schnorr signature scheme SDS is strongly unforgeable under definition 16.

258 Similarly, in the ROM we can trivially guarantee H is pseudorandom, whereas for our
259 setting we require an additional assumption:

▶ **Assumption 2.** The hash function generator $HGen$ is such that for all n

$$|\Pr[H \leftarrow HGen(n); (m_0, m_1) \leftarrow \mathcal{A}(H); b \leftarrow \{0,1\}; r \leftarrow \{0,1\}^\lambda : b = \mathcal{A}(H(m_b || r))] - 1/2|$$

260 is negligible in λ .

261 4 Ledger Model

262 The ledger model to which we tailor our light client design is a UTxO ledger with multi-asset
263 support ($UTxO_{ma}$), first introduced in [13]. The UTxO ledger model, such as the one used

264 by BitCoin [37], Ergo [22], and Cardano [10], maintains a record, called the *UTxO set*, of
 265 transaction outputs added by transactions that have been applied throughout its history, but
 266 not yet spent by subsequent transactions. We chose the UTxO_{ma} ledger because it allows
 267 us to demonstrate relevant usecases of our light client design (which would also work for a
 268 single-asset UTxO ledger), without introducing unnecessarily complexity of the Extended
 269 UTxO ledger. For completeness, and in order to establish notation, we include an overview
 270 of the UTxO_{ma} model. For additional notation explanation, see Figure 9.

271 **Blocks and Ledger States.** A block $\text{Block} = \text{Header} \times [\text{Tx}]$ is a data structure used to
 272 update the state of the ledger by applying a list of transactions $ltx \in [\text{Tx}]$ contained in the
 273 block, as well as doing some other checks and updates we do not model here. Among other
 274 data, the block header contains a slot number field $\text{slot} : \text{Header} \rightarrow \text{Slot}$, which represents the
 275 blockchain time at which the block is produced.

276 The ledger state is a data structure which is updated by applying blocks incoming on the
 277 network. The ledger state LState contains (among other data we do not model here) the UTxO
 278 set field, $\text{utxo} : \text{LState} \rightarrow \text{UTxO}$. It also contains the parameter $\text{minfee} : \text{LState} \in \mathbb{N}$, which is
 279 the minimum fee a transaction must pay. A block b can extend the blockchain (i.e. update
 280 the current state $s \in \text{LState}$) whenever the function

281 $\text{checkBlock} : \text{LState} \times \text{Block} \rightarrow \{0,1\}$

282 applied as $\text{checkBlock}(s, b)$ returns 1. Then, the updated block state is computed by
 283 $\text{updateState}(s, b)$. We do not give full specifications of checkBlock or updateState , as they are
 284 not required to model our light client approach.

285 **Full Nodes.** Let s_0 be some verified state, e.g. a genesis state, or a verified checkpoint state,
 286 and suppose $[b_0, \dots, b_k]$ is a list of blocks that have been disseminated across the network since
 287 the time slot of s_0 . We assume that a *full node* is one that is able to (1) compute the current
 288 state s_k by applying them in sequence, i.e. computing $s_{i+1} = \text{updateState}(s_i, b_i)$ for $0 \leq i < k$,
 289 (2) is able to determine if all of the blocks in the list are valid (i.e. $\text{checkBlock}(s_i, b_i) = 1$), and
 290 (3) can produce, send, and receive blocks on the network. If all the blocks are valid, we say
 291 that $[b_1, \dots, b_k]$ forms a valid blockchain.

292 **Ledger State and the UTxO Set.** The state of a UTxO-based ledger necessarily contains
 293 a UTxO set. While realistic ledgers often contain additional information in their state, in our
 294 model, the ledger state is just the UTxO set. The UTxO set is a finite map, $\text{UTxO} := \text{TxIn} \mapsto$
 295 TxOut . A transaction updates the UTxO set by either adding and removing entries.

296 **Transactions.** A transaction is the following data structure :

297 $\text{Tx} = (\text{inputs} : \text{Set TxIn}, \text{outputs} : [\text{TxOut}], \text{validityInterval} : \text{Interval}[\text{Slot}],$
 $\text{mint} : \text{Value}, \text{fee} : \mathbb{N}, \text{aux} : \mathbb{H}, \text{sigs} : \text{Signature})$

298 An input $(txid, ix) \in \text{TxIn} := \text{TxId} \times \mathbb{N}$ is a pair of a transaction ID and a natural number.
 299 When a transaction is applied to a UTxO set, its set of inputs is used to identify the entries
 300 which the transaction is removing from the set. In each input, $txid \in \text{TxId} := \mathbb{H}$ is the hash of
 301 a (previous) transaction that added that entry to the UTxO, and ix is the index of that output
 302 in the list of outputs of that transaction.

303 An output $(s, v) \in \text{TxOut} := \text{Script} \times \text{Value}$ is a pair of a script s which specifies some
 304 constraints that are checked when the output is spent, and the assets v contained in the output.
 305 Note here that our ledger model has native *multi-asset support*, following the scheme outlined
 306 in prior work [13]. That is, the $v \in \text{Value}$ in the output contains not only a quantity of a single
 307 currency (like BitCoin), but an arbitrary finite number of different types of assets with unique
 308 identifiers, alongside their quantities.

309 The list `outputs` of outputs of a transaction tx is used to construct a set of UTxO entries that
 310 will be added to the UTxO set, such that the unique identifier $txin$ of each output $o \in \text{outputs } tx$
 311 consists of the transaction hash `txid` tx , and the index of o in the list `outputs` tx . The entries
 312 added to the UTxO set by tx are computed in this way by `mkOuts`, see Figure 10.

313 The interval `validityInterval` specifies the range of slot numbers for which a transaction can
 314 be valid, or `nothing` whenever there is no limitation. The field `sigs`:`Signature`:=`PubKey`→ \mathbb{H}
 315 is a set of public keys, associated with their signatures on the the transaction (excluding `sigs`
 316 itself). The `fee` is the amount of primary currency a transaction pays as a system fee, which
 317 is checked to be at least the required fee `minfee`.

318 The `mint` field represents the assets being minted or burned by the transaction. Assets with
 319 positive quantities are said to be minted, while those with negative quantities are burned. When
 320 a transaction is applied, the constraints specified by every $p \in \text{Policy} := \text{Script}$ of each type of
 321 asset specified in this field are checked to make sure minting/burning of this type and quantity of
 322 asset is allowed. The `aux` is a field for arbitrary extra data encoded as a bytestring, e.g. a "note".

323 **Ledger State Update.** Given a UTxO set $utxo$ and a transaction tx , the function
 324 `updateUTxO`: $\text{UTxO} \times \text{Tx} \rightarrow \text{UTxO}$ computes the updated UTxO set by adding and removing
 325 the appropriate entries :

$$326 \quad \text{updateUTxO } (utxo, tx) = \{ i \mapsto o \in utxo \mid i \notin \text{inputs } (tx) \} \cup \text{mkOuts}(tx)$$

327 While an update to the UTxO set can be computed for any transaction, only transactions
 328 that are *valid* for a given set are allowed to perform an update to the ledger state. For a given
 329 $utxo$ set, a transaction tx is valid whenever the function $\text{checkTx} : (\text{Slot} \times \mathbb{N}) \times \text{UTxO} \times \text{Tx} \rightarrow \{0,1\}$,
 330 applied as $\text{checkTx } ((slot, fee), utxo, tx)$, returns 1. The `checkTx` function is the conjunction
 331 of the constraints specified in Section B.1. For a given block b and state s , the function
 332 `updateState` (s, b) updates s with the list of transactions $\pi_2 b = [tx_1; \dots; tx_k]$ in such a way that the
 333 update to the UTxO set contained in s is computed by applying the transactions in sequence, i.e.

$$334 \quad \text{updateUTxO } (\text{updateUTxO } ((\dots utxo s\dots), tx_{k-1}), tx_k) = \text{utxo } (\text{updateState } (s, b))$$

335 For each transaction in the list, $\text{checkTx } ((slot b, fee s), utxo_i, tx_i) = 1$ is first checked, and
 336 the entire block is considered invalid if this check fails. As part of checking transaction validity,
 337 constraints of every `Script` run by the transaction are checked. The constructors and evaluation
 338 of `Script` is given in Figure 11, and `MOf` is given in 10. A script, for a given set of signer keys
 339 khs and slot numbers $s1, s2$, can be defined to check that (some specific) m of them have signed
 340 the transaction (are included in the domain of `Signature`), and/or that the validity interval of
 341 the transaction starts after $s1$ and/or ends before $s2$.

342 5 Light Client Specification

343 We model light clients in terms of their functionality, constraints, and protocols for communica-
 344 tion with full nodes and potential service providers. By functionality, we refer to the interaction
 345 between a light client and its user(s). We describe the functionality of a light client in terms of
 346 constructing (posting) and verifying resolved *intents*. A light client specifies the intent it wants
 347 resolved using a special domain-specific language (DSL). Then, the client sends the intent to a
 348 service provider (SP), who is incentivized to respond to this intent with an *abstract transaction*.
 349 An abstract transaction is a transaction that has been modified to partially conceal data in
 350 a way that (i) allows the light client to verify that their intent has been resolved in the original
 351 transaction and (ii) does not allow the client to use the abstract transaction to form a valid
 352 transaction that avoids paying the service provider fee.

```

TxAbs = (outputs:[TxOut], validityInterval:Interval[Slot],
          mint:Value, fee:N, sigKeys:Set PubKey)

mkAbs : Tx → TxAbs
mkAbs tx = tx { outputs=tx.outputs,
                  outputs[0].Value=-1, //Leftover balance to change address
                  validityInterval=tx.validityInterval,
                  mint=tx.mint, fee=tx.fee
                  sigKeys=dom (tx.sigs) }

```

■ **Figure 1** Abstract transaction type `TxAbs` and constructor `mkAbs`

353 A light client's constraints may be on the communication cost to answer a series of queries,
 354 available state/data storage, and assumptions about connectivity. As a consequence, a par-
 355 ticular light client design may be limited in the intents it can construct even if a broader class
 356 of intents is supported by available software.

357 5.1 Intent specification

358 An *abstract transaction*, `TxAbs`, is the data structure that a light client receives from the service
 359 provider instead of the full plaintext transaction (see Fig. 1). In `TxAbs`, the `aux` and the `inputs`
 360 fields are removed, and instead of the `Signature` field, it contains only the signing keys `sigKeys`
 361 and not the corresponding signatures. The value of the first output is zeroed out to allow for
 362 any leftover input funds to be returned to the LC's change address. The intent a light client
 363 constructs is any expression in the DSL $\mathcal{I}_{\text{post}}$, see Figure 2.

364 For the implementation of the example usecases of our design, we require the definition
 365 of the following functions:

- 366 (i) The function that constructs a transaction based on the specification (this function may fail,
 367 outputting nothing instead of a transaction, if the intent cannot be resolved), see Fig. 12.

368 $\text{mkToSpec} : \text{LState} \times \mathcal{I}_{\text{post}} \rightarrow \text{Tx}^?$

- 369 (ii) The function that checks that a given abstract transaction matches the intent specification,
 370 see Figure 2

371 $\text{chkSpec} (\mathcal{I}_{\text{post}} \times \text{TxAbs}) \rightarrow \{0,1\}$

372 Note that the function `mkToSpec` constructs a transaction by editing an initial transaction
 373 $\text{initTx}_{a,mf}$ (see Figure 12). This transaction includes a minimum fee mf as well as a tip tip to
 374 the service provider, output to its address a . These functions are defined such that for any l, i ,
 375 where $\text{mkToSpec}(l, i) \neq \text{nothing}$, necessarily $\llbracket i \rrbracket_{\text{DSL}}(\text{mkAbs}(\text{mkToSpec}(l, i))) = 1$. The design
 376 of the DSL is minimal, and is meant only to showcase the operation of the light client protocol.
 377 To expand the functionality of the light client, more constraints need to be included in the
 378 DSL, e.g., a way to limit transaction size, or support for specifying desired token exchanges.
 379 We leave this for future work.

380 We showcase two examples in the following, for a given ledger state l :

- 381 (txi₁) Intent to mint some token t within an interval of length j and maximum fee $f \leq \text{minfee } l$
 382 (note that this will not be a valid transaction since it will have no inputs):

383 $i_1 = \text{AndExps} [\text{MustMint } t; \text{MaxInterval } j; \text{MaxFee } f; \text{ChangeTo } s]$

23:10 Cavefish

$\mathcal{I}_{\text{post}}$ CONSTRUCTORS

MustMint	: Value → $\mathcal{I}_{\text{post}}$
SpendFrom	: Script → $\mathcal{I}_{\text{post}}$
MaxInterval	: Slot → $\mathcal{I}_{\text{post}}$
PayTo	: (Value × Script) → $\mathcal{I}_{\text{post}}$
ChangeTo	: Script → $\mathcal{I}_{\text{post}}$
MaxFee	: $\mathbb{N} \rightarrow \mathcal{I}_{\text{post}}$
AndExps	: $[\mathcal{I}_{\text{post}}] \rightarrow \mathcal{I}_{\text{post}}$

EVALUATION OF $\mathcal{I}_{\text{post}}$

$\llbracket _ \rrbracket_{\text{DSL}}$: $\mathcal{I}_{\text{post}} \rightarrow \text{TxAbs} \rightarrow \{0,1\}$
$\llbracket \text{MustMint } v \rrbracket_{\text{DSL}}(tx)$	= $v \leq tx.\text{mint}$
$\llbracket \text{SpendFrom } s \rrbracket_{\text{DSL}}(tx)$	= $\llbracket s \rrbracket(\text{dom}(tx.\text{sigs}), tx.\text{validityInterval})$
$\llbracket \text{MaxInterval } i \rrbracket_{\text{DSL}}(tx)$	= $(tx.\text{validityInterval})_2 - (tx.\text{validityInterval})_1 \leq i$
$\llbracket \text{PayTo } (s,v) \rrbracket_{\text{DSL}}(tx)$	= $(s,v) \in tx.\text{outputs}$
$\llbracket \text{ChangeTo } s \rrbracket_{\text{DSL}}(tx)$	= $(s, \text{consumed} - \text{produced}) \in tx.\text{outputs}$
$\llbracket \text{MaxFee } f \rrbracket_{\text{DSL}}(tx)$	= $tx.\text{fee} \leq f$
$\llbracket \text{AndExps } [a1;a2;...;ak] \rrbracket_{\text{DSL}}(tx)$	= $(\llbracket a1 \rrbracket_{\text{DSL}} tx) \wedge (\llbracket a2 \rrbracket_{\text{DSL}} tx) \wedge \dots \wedge (\llbracket ak \rrbracket_{\text{DSL}} tx)$

Figure 2 $\mathcal{I}_{\text{post}}$ constructors and evaluation

- 36 $\exists (tx_i)_2$) Intent to pay x from outputs locked by RequireSig $k1$ to RequireSig $k2$ (note that this intent
385 may not be possible to resolve in the key $k1$ does not have sufficient funds on the ledger l)
- 386 $i_1 = \text{AndExps} [\text{SpendFrom}(\text{RequireSig } k1); \text{PayTo}(\text{RequireSig } k2, x); \text{ChangeTo}(\text{RequireSig } k1)]$
- 387 The resulting transactions are given in Figure 13, assuming the intents can be resolved. The
388 minimum required capacity of our light client is to perform all the actions that are specified
389 in the protocol we describe in the next section.

390 6 Light Client Protocols

- 391 We give an overview on our light client protocol and its components. Our protocol maximizes
392 space and communication efficiency by essentially pushing all storage and lookup tasks to
393 the SP, whilst guaranteeing the resulting transaction is safe for the light client to sign, and
394 compensating their SP for their effort. This is non-trivial as the transaction must be signed
395 sight-unseen (blinded) otherwise, the light client could simply remove the payment to the SP
396 and submit the modified transaction.

397 To satisfy both parties, we rely on cryptographic tools as well as some established features
398 in implemented blockchains such as transaction validity periods and handling of invalid transac-
399 tions. In the rest of this section, we make the following assumption: all communication between
400 a light client and a service provider happens over a secure channel (e.g. a TLS connection).

401 6.1 Definitions and requirements

- 402 ▶ **Definition 2** (Correctness). A transaction building protocol is **correct** if for an honest light
403 client, and service provider and any state $LState_i \in LState$ and intent $int_{\text{post}} \in \mathcal{I}_{\text{post}}$ such that:

404 if $\text{mkToSpec}(\text{LState}_i, \text{int}_{\text{post}}) \neq \perp$, then the protocol completes and the transaction Tx' produced
 405 by the SP is such that $\text{checkTxL}(\text{LState}_i, \text{Tx}') = 1$.

406 We use $\text{checkTxL}(l, \text{tx})$ as shorthand for $\text{checkTx}((\text{slot}, \text{fee}), \text{utxo}, \text{tx})$, with ledger state l
 407 containing UTxO state utxo and current slot number slot and minimum fee fee .

408 ▶ **Definition 3 (Safety).** A transaction building protocol is *safe* if for an honest light client
 409 and any service provider SP we have that: if transaction Tx' is produced by the SP is such that
 410 $\text{checkTxL}(\text{LState}, \text{Tx}') = 1$, then it must be that $\text{chkSpec}(\text{int}_{\text{post}}, \text{Tx}') = 1$.

411 The safety property is derived from two factors: first, the unforgeability property of
 412 WBPS[P] will ensure that only transactions that satisfy the given predicate will be signed.
 413 Second, our assumptions on ledger rules ensure that signed transactions involving spent UTXOs
 414 will have no effect (as opposed to something undesirable to the user).

415 ▶ **Definition 4 (Private until posted).** A transaction building protocol has the *private until*
 416 *posted* (PuP) property if a PPT adversarial client $C_{\mathcal{A}}$ cannot win the following experiment
 417 with probability significantly higher than $\frac{1}{2}$.

418 For an honest service provider SP and any client $C_{\mathcal{A}}$ we have that: client provides
 419 2 states $\text{LState}_0, \text{LState}_1$ and one intent $\text{int}_{\text{post}} \in \mathcal{I}_{\text{post}}$ such that for $i=0,1$ it is
 420 $\text{mkToSpec}(\text{LState}_i, \text{int}_{\text{post}}) \neq \perp$. The experiment flips a coin d and runs the protocol between SP and C^* using state LState_d .
 421 $C_{\mathcal{A}}$ outputs d^* . The adversarial client wins iff $d=d^*$.

422 The Private until posted property of our protocol is derived from Theorem 9 and the
 423 protocol structure. Given an adversary against \mathcal{A} PuP, we can build an adversary against
 424 the weak blindness of the WBPS[P] scheme by running the SP once on LState_0 and once on
 425 LState_1 to produce two challenge messages for the blindness security game and simulate the
 426 role of the challenge towards \mathcal{A} by passing on the rest of the protocol messages (the experiment
 427 enforces random padding internally). If \mathcal{A} can distinguish which ledger was used in producing
 428 the messages we can use the same index as our guess in the PuP game.

429 6.2 Cavefish: a communication-optimal transaction building protocol

430 We give an overview on our light client protocol and its components. As the protocol is
 431 uniquely suited for UTxO-based blockchains, we develop our approach to support the general
 432 ledger model explained in Section 4 and the intents set forth in Section 5.1.

433 6.2.1 Security model and requirements

434 **Rational Actors.** For liveness, we assume that protocol participants are rational in the sense
 435 that they will opt to complete tasks that benefit them directly (i.e. in the form of rewards, or in the
 436 form of a desired transaction being posted). For safety we assume parties can be fully malicious.

437 **Transaction Expiry.** For safety, we require that transactions can be set to be valid only
 438 within a certain time window. This is to prevent trivial attacks where a malicious SP delays
 439 posting a transaction indefinitely. The user could contact a different SP to create a second
 440 transaction with similar parameters (e.g. by paying the same recipient with different coins)
 441 with the risk that the first SP would also post their transaction, making the user pay twice.

442 **Idempotent Doublespends.** For safety we also require that attempted double spends are
 443 idempotent, i.e. they are simply not valid transactions and do not penalize (slash) the user
 444 for signing them.

445 **Trusted Setup.** We specify a proof system using a trapdoor, which is a common characteristic
 446 of efficient options such as Groth16 [27] or Plonk [25]. It is possible to use alternatives featuring

447 transparent setup with some impact to communication size. In the case of a transparent setup
 448 for NArg, we still need to generate keys for PKE. In the case of Elgamal in the random oracle
 449 model, this can be accomplished by hashing an unpredictable string into a group element.

450 6.3 Protocol description

451 The protocol operates as follows: the light client connects to the SP and sends a posting intent
 452 int_{post} , as well as a public key that has received payments (and, optionally a BIP-032 chaincode
 453 so that the SP can also derive child keys from the initial one –we explain the details in Section
 454 8). The SP will then search for UTXOs belonging to the client and form a transaction that
 455 satisfies the intent int_{post} , including in the transaction a tip for the SP’s own work. Lastly,
 456 the SP adds a random note aux_{nt} to the note field of the transaction. Then, the client and SP
 457 engage in a weakly blind predicate signature protocol so as to efficiently have the client sign the
 458 transaction whilst ensuring that (1) the client needs to sign before learning the full transaction
 459 (typically the used UTXOs are hidden), and (2) that the signed transaction satisfies int_{post}
 460 by means of satisfying the corresponding predicate.

461 We can instantiate our protocol for two kinds of intents: “light” and “extended”. In the
 462 light version, we support intents where the transaction abstract does not give the light client
 463 any information the SP does not wish to disclose—by design, the abstract hides the input
 464 transactions, but it could reveal, e.g. some state details of a smart contract.

465 A practical example of such intents is $int_{\text{light}}(\text{tx}) := \{\text{return } (\text{makeAbs}(\text{tx}) = \text{makeAbs}(\text{tx}_0))\}$
 466 i.e. the light client fully determines the requested transaction (including SP fee) to be tx_0 ,
 467 apart from the input UTXOs (which it does not know) and the note field aux which is free
 468 to the SP to set. This implies that TxAbs provides no new information to the light client.
 469 This reduces the computational effort by SP invested in running the NArg prover as chkSpec
 470 consists only of substring equality checks. We note here that no range check is necessary for
 471 values: makeAbs (Section 5.1) retains the prescribed value assigned to all outputs except the
 472 first (which is used as a “change” address), i.e., the only option for the SP is to assign all excess
 473 value to the change address.

474 Alternatively, in the extended variant, we set $\text{makeAbs}(x) := \{\text{return 1}\}$, and change the
 475 predicate check from $\text{chkSpec } (\mathcal{I}_{\text{post}} \times \text{TxAbs}) \rightarrow \{0,1\}$ to $\text{chkSpec } (\mathcal{I}_{\text{post}} \times \text{Tx}) \rightarrow \{0,1\}$. This too
 476 ensures that no information is leaked by TxAbs . Our results hold for both variants.

477 We show our complete light client protocol in Figure 3, leveraging the weakly blind signa-
 478 tures predicate of section 7.2. For simplicity, the figure describes the process for a single UTXO
 479 signature, if more are needed, multiple WBPSsessions can be run sharing the same com_{tx} .

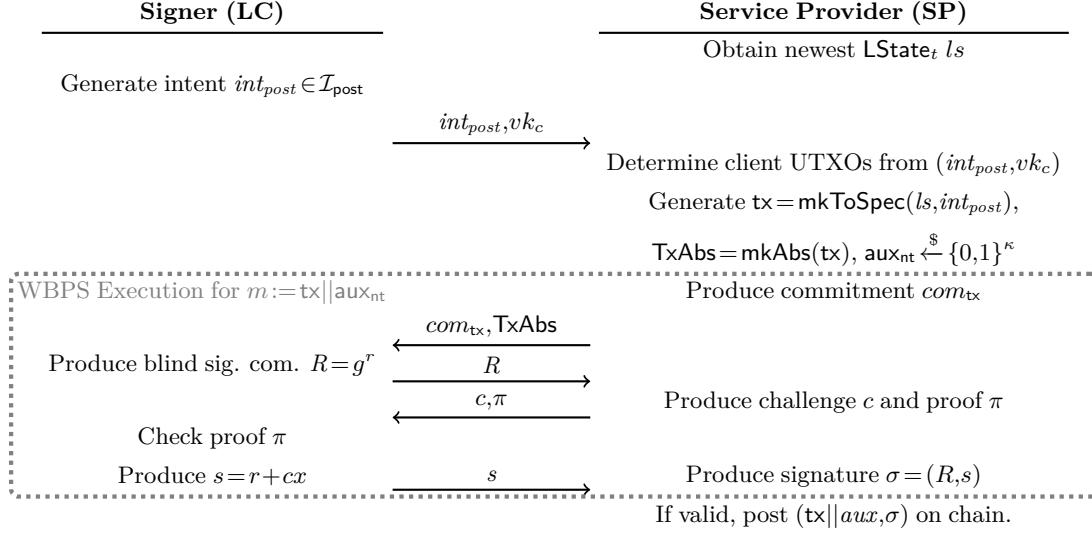
480 6.3.1 Security

481 ▶ **Theorem 5.** *If the ledger rules include the Idempotent Doublespends property and WBPS[P]
 482 is unforgeable, the protocol of Section 6.3 is safe.*

483 **Proof.** The safety property is derived from two factors: first, the unforgeability property of
 484 WBPS[P] (Theorem 10) will ensure that only transactions that satisfy the given predicate will
 485 be signed. Second, our assumptions on ledger rules ensure that signed transactions involving
 486 spent UTXOs will have no effect (as opposed to something undesirable to the user). ◀

487 ▶ **Theorem 6.** *If WBPS[P] scheme is weakly blind, the protocol of Section 6.3 is private until
 488 posted.*

489 **Proof.** The Private until posted property of our protocol is derived from Theorem 9 and the
 490 protocol structure. Given an adversary against \mathcal{A} PuP, we can build an adversary against



■ **Figure 3** Cavefish protocol featuring WBPS scheme. A detailed description of the WBPS scheme is in Fig. 6.

491 the weak blindness of the WBPS[P] scheme by running the SP once on LState_0 and once on
 492 LState_1 to produce two challenge messages for the blindness security game and simulate the
 493 role of the challenge towards \mathcal{A} by passing on the rest of the protocol messages (the experiment
 494 enforces random padding internally). If A can distinguish which ledger was used in producing
 495 the messages we can use the same index as our guess in the PuP game. ◀

496 6.4 Discussion

497 **Suboptimal Transactions.** In response to an LC query, it is possible for an SP to construct
 498 a *sub-optimal* transaction (in terms of cost) which nevertheless matches the LC’s specification.
 499 For example, the LC may include a larger-than-necessary system fee in the transaction, or not
 500 provide the LC with the best available exchange price for a specific token. We consider cost
 501 optimizations an orthogonal issue that can be addressed by e.g. a market system.

502 **Liveness against Malicious SPs.** Malicious SPs are able to force delays for the user by
 503 completing the signing and never posting the transaction. This forces the user to wait until the
 504 signed transaction is no longer valid before retrying. Otherwise, as intents are not necessarily
 505 uniquely satisfied, it is possible that the first (malicious) node may be able to post the first
 506 transaction (which was purposefully delayed) after the posting of the second one. This can
 507 cause the user to e.g. double pay for a service as well as the full node fee. Even so, given enough
 508 retries the user will reach an honest SP and the transaction will be posted.

509 6.4.1 Multi-SP protocols and optimality

510 However, we can also formulate a version of this protocol where the LC instead sends the
 511 specification to multiple SPs, selecting the best response, and engaging in the rest of the
 512 protocol only with that SP.

513 To do this, let $\text{opt}: \text{TxAbs} \rightarrow \mathbb{Z}$ be a function that rates abstract transactions to express LC’s
 514 preferences. For example, the total amount of primary tokens spent by the transaction can

515 be such a function : $\text{opt } \text{tx} = \text{coinValue} (\sum_{o \in \text{txd.spentOuts}} o.\text{value})$. A transaction tx is *optimal*
 516 in a set $S \in \text{Set TxAbs}$ when $\text{opt } \text{tx} = \min \{ \text{opt } \text{tx}' \mid \text{tx}' \in S \}$.

517 To get a multi-SP protocol, the first step of the single-SP protocol in Figure 3 must be
 518 augmented, so that the pair $(\text{opt}, \text{int}_{\text{post}}$ is sent to each SP instead of just sending int_{post} . Upon
 519 receiving transactions $\text{tx}_{A,i}$, with $0 \leq i < k$, from each of the k SPs responding to LC's query,
 520 LC will engage in the rest of the protocol only with the sender of $\text{tx}_{A,i}$.

521 Note that the opt function and the specification serve different purposes in the protocol. The
 522 specification is checked, and any response transaction that does not satisfy it is discarded. On
 523 the other hand, it is not required that a transaction be optimal across all possible specification-
 524 satisfying transactions. The kind of sophisticated optimization (such as what is required, e.g.,
 525 for optimized order-matching) would require an entirely distinct set of tools for demonstrating
 526 the optimality result, such as ZK proofs about the full blockchain state (rather than just the
 527 associated transaction), together with evidence that the proof is about state that is *sufficiently*
 528 *current* (see Section 6.4). For example, a proof that there were no better offers for a
 529 specific token available on the ledger at the time the SP produced a response to the LC. We
 530 do not assume that either the SP or the LC are necessarily capable of performing or verifying
 531 (resp.) optimality according to the function LC requested be optimized, but an LC is capable
 532 to comparing transactions using opt .

533 7 Blind Signatures for Abstract Transactions

534 In order to allow the light to sign an abstract transaction we implement blind signatures
 535 on (partially) blinded transaction objects. At the minimum the SP hides the inputs to the
 536 transaction, i.e., the references to UTxO objects which are present in the ledger and required to
 537 cover the transaction. This ensures that the light client cannot simply obtain the transaction
 538 from the SP, then edit it to remove the SPs compensation.

539 Our construction is inspired by the predicate blind signature mechanism in [24]. It realizes
 540 a concurrently secure blind and partially blind signing protocol resulting in standard Schnorr
 541 signatures. This is a core property as it allows our protocol to be compatible with popular
 542 blockchains such as Bitcoin and Cardano without any modifications.

543 Given our application, the unlinkability property of blind signatures is not of much benefit:
 544 signed transactions are somewhat infrequent, often have different payees or payment sums (i.e.
 545 different predicates) and are indirectly timestamped due to their time of posting on the blockchain.
 546 Further, there seems to be little benefit in preventing the light client from knowing which protocol
 547 section produced which payment. For this reason, we do away with the requirement of the signer
 548 being unable to distinguish between messages signed on different sessions. We only require
 549 that the signer is unable to distinguish the message *before* the signature-message pair is posted.

550 This weakening of the blindness property brings about the benefit that we can remove the
 551 blinding operations from the protocol (for a small efficiency increase) and also from the proof of
 552 correctness of the challenge generation (where the gain is more significant as it removes group
 553 exponentiations using foreign field arithmetic). The resulting scheme is similar in operation
 554 to the Signatures over Blocks of Committed Messages construction of Bobolz et. al. [7] with
 555 the addition of predicate checking.

556 7.1 Weakly blind predicate signatures

557 We adapt the definitions of [24] to account for the weak variant of blindness.

558 A WBPS scheme is parameterized by a family of polynomial-time-computable predicates,
 559 which are implemented by a p.t. algorithm P , the predicate compiler: on input a predicate

560 description $prd \in \{0,1\}^*$ and a message $m \in \{0,1\}^*$, P returns 1 or 0 indicating whether m
 561 satisfies prd .

562 A WBPS scheme $\text{WBPS}[P]$ for predicate P is defined by the following algorithms. We focus
 563 on schemes with 2-round (i.e., 4-message) signing protocols for concreteness.

- 564 ■ $\text{Setup}(1^\lambda) \rightarrow par$: the setup algorithm, on input the security parameter, outputs public
 565 parameters par , which define a message space \mathcal{M}_{par} .
- 566 ■ $\text{KeyGen}(par) \rightarrow (sk, vk)$: the key generation algorithm, on input the parameters par , out-
 567 puts a signing/verification key pair (sk, vk) , which implicitly contain par , i.e., $vk = (par, vk')$.
 568 —
- 569 ■ $\langle \text{Sign}(sk, prd), \text{User}(vk, prd, m) \rangle \rightarrow (b, \sigma)$: an interactive protocol with shared input par
 570 (implicit in sk and vk) and a predicate prd is run between the signer and user. The signer
 571 takes a secret key sk as private input, the user's private input is a verification key vk and a
 572 message m . The signer outputs $d=1$ if the interaction succeeds and $d=0$ otherwise, while
 573 the user outputs a signature σ if it succeeds, and \perp otherwise.
- 574 ■ $\text{Ver}(vk, m, \sigma) \rightarrow 0/1$: the (deterministic) verification algorithm, on input a verification key
 575 vk , a message m and a signature σ , outputs 1 if σ is valid on m under vk and 0 otherwise.

576 For a 2-round protocol the interaction $\langle \text{Sign}(sk, prd), \text{User}(vk, prd, m) \rangle \rightarrow (d, \sigma)$ can be
 577 realized by the following algorithms:

$$\begin{aligned} 578 \quad (txt_{U,0}, st_{U,0}) &\leftarrow \text{User}_0(vk, prd, m) & (txt_{S,1}, st_S) &\leftarrow \text{Sign}_1(sk, prd, txt_{U,0}) \\ 579 \quad (txt_{U,1}, st_{U,1}) &\leftarrow \text{User}_1(st_{U,0}, txt_{S,1}) & (txt_{S,2}, d) &\leftarrow \text{Sign}_2(st_S, txt_{U,1}) \\ 580 \quad \sigma &\leftarrow \text{User}_2(st_{U,1}, txt_{S,2}) \end{aligned}$$

581 We write $(d, \sigma) \leftarrow \langle \text{Sign}(sk, prd), \text{User}(vk, prd, m) \rangle$ as shorthand for the above sequence.

582 ▶ **Definition 7.** A WBPS scheme $\text{WBPS}[P]$ satisfies weak blindness if for all p.p.t. adversaries
 583 \mathcal{A} :

$$Adv_{\text{WBPS}[P], \mathcal{A}}^{\text{BLD}}(\lambda) := \Pr[\text{BLD}_{\text{WBPS}[P]}^{A,1}(\lambda) - \text{BLD}_{\text{WBPS}[P]}^{A,0}(\lambda)] \text{ is negligible in } \lambda.$$

584 We note that the BLD experiment mauls the message by appending a random bitstring r of
 585 length λ . This is appropriate to our setting, where the SP is explicitly allowed to use the “notes”
 586 field of a transaction to add a randomizer. Without this mauling, any indistinguishability-based
 587 definition would fail.

▶ **Definition 8.** . A WBPS scheme $\text{WBPS}[P]$ satisfies unforgeability if for all p.p.t. adversaries
 \mathcal{A}

$$Adv_{\text{WBPS}[P], \mathcal{A}}^{\text{EUF-CMA}}(1^\lambda) := \Pr[\text{CMA}_{\text{WBPS}[P]}^{\mathcal{A}}(1^\lambda) = 1] \text{ is negligible in } \lambda.$$

588 7.2 Cavefish scheme

589 Our Scheme operates as follows: given an intent (i.e. predicate predicate description)
 590 $prd := int_{post}$ and predicate compiler P s.t. $(P(prd))(tx) := \text{chkSpec}(int_{post}, \text{mkAbs}(tx))$ the
 591 requestor (i.e. the SP) encrypts the message, in our case the full transaction $m := tx$ and sends
 592 the commitment C to the signer, who returns a random group element R used in the final
 593 Schnorr signature. The requestor replies with the Schnorr challenge c as well as a proof π of
 594 its correct construction, and of the fact that tx satisfies int_{post} .

595 We instantiate the scheme with a parametrisable NArg for the following relation,

$$596 \quad R_{\text{Cavefish}}(\underbrace{(q, \mathbb{G}, G, H)}_{\text{parameters } par}, \overbrace{(X, R, com_{tx}, TxAbs, c, int_{post})}^{\text{known statement } \theta}, \underbrace{(tx || aux_{nt}, \rho)}_{\text{witness } \omega}) \quad (1)$$

$\text{BLD}_{\text{WBPS}[P]}^{A,b}(\lambda)$	$\text{ChalUser}(msg = \emptyset)$
$par \leftarrow \text{Setup}(1^\lambda)$	if $sess = \text{await}$
$(m_0, m_1, prd, vk', st_A) \leftarrow \mathcal{A}_1(par)$	$sess \leftarrow \text{closed}$
if $P(prd, m_0) = 0$ or $P(prd, m_1) = 0$	$\sigma \leftarrow User_2(st_u, msg)$
then return 0	if $sess = \text{closed}$
$\text{aux}_{\text{nt}} \leftarrow \{0,1\}^\lambda$	$msg' \leftarrow Ver(vk, m, \sigma)$
$m \leftarrow m_b \text{aux}_{\text{nt}}$	if $sess = \text{open}$
$vk \leftarrow (par, vk')$	$sess \leftarrow \text{await}$
$sess \leftarrow \text{init}$	$(msg', st_u) \leftarrow User_1(st_u, msg)$
$b^* \leftarrow \mathcal{A}_2^{\text{ChalUser}}(st_A)$	if $sess = \text{init}$
return $b = b^*$	$sess \leftarrow \text{open}$
	$(msg', st_u) \leftarrow User_0(vk, prd, m)$
	return msg'

Figure 4 Weak Blindness experiment $\text{BLD}_{\text{WBPS}[P]}^{A,b}(\lambda)$ for a WBPS scheme with predicate compiler P , adversary \mathcal{A} and parameter b .

$\text{CMA}_{\text{WBPS}[P]}^A(\lambda)$	$\text{SigInit}(prd, msg_{U,0})$
$par \leftarrow \text{Setup}(1^\lambda)$	$S \leftarrow S+1; P \leftarrow P+1$
$(sk, vk) \leftarrow \text{KeyGen}(par)$	$(msg, st_S) \leftarrow \text{Sign}_1(sk, prd, msg_{U,0})$
$Q \leftarrow 0$	$prds \leftarrow prd$
$S \leftarrow 0$	return msg
$P \leftarrow 0$	$\text{SigComplete}(s, msg_{U,1})$
$(\vec{m}^*, \vec{\sigma}^*, \vec{prd}^*) \leftarrow \mathcal{A}^{\text{SigInit}, \text{SigComplete}}(vk)$	if $s > S$ or $st_s = \perp$
$n \leftarrow \vec{m}^* $	then return \perp
if $\prod_i Ver(vk, m_i^*, \sigma_i^*) \neq 1$ return 0	$(msg, d) \leftarrow \text{Sign}_2(st_s, msg_{U,1})$
if $\prod_i prd_i^*(m_i^*) \neq 1$ return 0	if $d = 1$ then $Q \leftarrow Q+1$
if $\exists i, j : (i \neq j) \wedge (m_i^*, \sigma_i^*) = (m_j^*, \sigma_j^*)$ return 0	$P \leftarrow P-1$
if $n > Q$ return 1	$st_S \leftarrow \perp$
if $\#\rho \in \text{Perm}(S)$:	return msg
$\forall i \leq n : (prd_{\rho(i)}(m_i) = 1) \wedge (st_{\rho(i)} = \perp)$	
then return 1	
return 0	

Figure 5 Chosen Message Unforgeability Experiment $\text{CMA}_{\text{WBPS}[P]}^A(\lambda)$ for a WBPS scheme with predicate compiler P , adversary \mathcal{A}

597 checking that $\text{chkSpec}(int_{post}, \text{TxAbs}) = 1 \wedge m = \text{tx} || \text{aux}_{\text{nt}} \wedge C = \text{PKE}.\text{Enc}(m; \rho) \wedge c = \text{H}(R, X, m) \wedge$
598 $\text{TxAbs} = \text{mkAbs}(\text{tx})$. For the extended version, we alter the predicate compiler to $(P_{\text{ext}}$ such
599 that $(P_{\text{ext}}(prd))(\text{tx}) := \text{chkSpec}(int_{post}, \text{tx})$ (i.e. we check tx rather than $\text{mkAbs}(\text{tx})$).

600 The signer completes the signing only if the proof verifies correctly. We present the full
601 protocol in Figure 6.

602 7.3 Security

603 ► **Theorem 9.** *The WBPS[P] scheme of Figure 6 achieves weak blindness when the encryption
604 scheme PKE is IND-CPA secure, NArg has the zero knowledge property, assumption 2 holds
605 w.r.t. the hash generator HGen.*

606 **Proof.** We structure our proof as a sequence of games so that (1) two consecutive games
607 are computationally indistinguishable to the adversary and (2) in the final game, the view
608 of the adversary is independent of the challenge b . We define the initial game G_0 to be

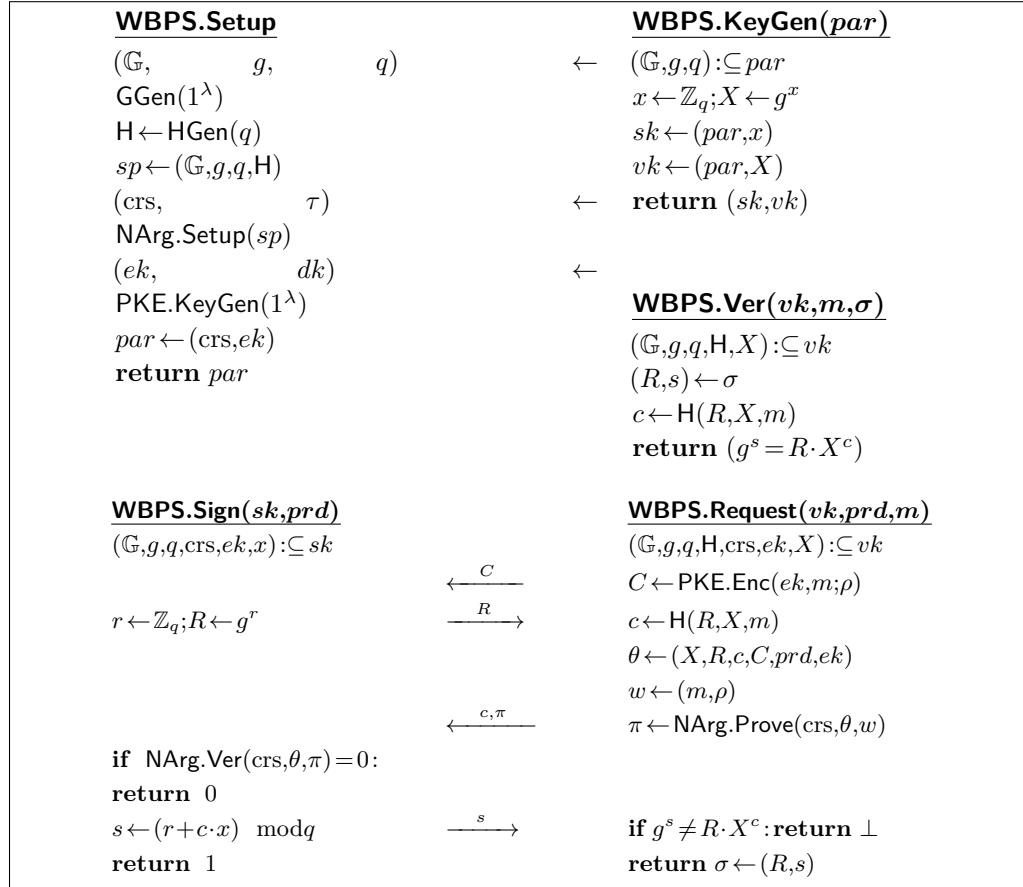


Figure 6 The weakly blind predicate Schnorr signature scheme WBPS[P]

609 $G_0 := \text{BLD}_{\text{WBPS}[P]}^{\mathcal{A},b}(\lambda)$. For G_1 we replace proofs with simulated ones. This is not detectable
610 by the adversary due to the zero knowledge property of **NArg**. Second, for game G_2 we replace
611 C with an encryption of a fixed message \bar{m} rather than m_b , this too is undetectable⁷ due to
612 the IND-CPA security of **PKE**. For game G_3 we replace the challenge c to also use \bar{m} rather
613 than m_b in the hash. This is computationally indistinguishable due to assumption 2. \blacktriangleleft

614 ▶ **Theorem 10.** *The WBPS[P] scheme of Figure 6 achieves unforgeability when the encryption
615 scheme PKE is IND-CPA secure, NArg is sound, and assumption 1 holds w.r.t. the hash
616 generator HGen.*

617 **Proof.** We follow the proof of [24] with little changes. As before, we structure our proof as
618 a sequence of games so that the probability of adversarial success changes negligibly from one
619 game to the next. We define the initial game G_0 to be $G_0 := \text{CMA}_{\text{WBPS}[P]}^{\mathcal{A}}(\lambda)$.

620 For game G_1 , the experiment holds the decryption key for **PKE**, so that we extract the
621 message from C ahead of time, and thus predict the c value once the value of R has been
622 determined. If the value sent by the adversary is different from the predicted one, and the
623 accompanying proof π verifies we abort early, diverging from the original game. This only

⁷ This will also change the input of the simulator, but this is not an obstacle: we simply consider the simulator to be part of the adversary we are constructing against the IND-CPA security of **PKE**.

624 happens with negligible probability though, due to the adaptive soundness of \mathbf{NArg} (if π does
 625 not verify, G_0 would abort as well).

626 For game G_2 , we check to see if any of the adversarial signatures uses a message that
 627 belongs to a session that was not closed. If so, we abort early, diverging from G_1 . However,
 628 this implies that the second part of the signature s along with the secret key x can be used
 629 to calculate the discrete log of R . Thus, if the early abort occurs (i.e. adversary manages
 630 to win using unfinished sessions) with non negligible probability we can build a discrete log
 631 calculator by embedding a discrete log challenge in one of the R values. Otherwise, we can
 632 assume that almost all of the adversarial wins use at least one “new” message, and thus the
 633 success probability for G_1 and G_2 differ only negligibly .

634 For G_3 we build a reduction to the EUF-CMA security of Schnorr (Assumption 1). We
 635 remove key generation from the challenger, and instead obtain a vk from the Schnorr EUF-CMA
 636 security game. When we decrypt a message m , we query it on our signing oracle, obtaining
 637 R,s as the signature. The reduction forwards R to the adversary anticipating c,π . The checks
 638 introduced in G_1 ensure the c value matches the predicted one, and thus s correctly completes
 639 the signature. This ensures the simulation can complete the game. If the adversary wins the
 640 security game $\mathbf{WBPS}[P]$ security game, there must be a signature on a new message that the
 641 reduction uses to also win the Schnorr EUF-CMA game with the same probability. ◀

642 8 Analysis & Applications

643 **Storage and Communications Requirements.** We point out that our protocol is optimal
 644 in terms of light client storage and light client & SP communication. The light client (LC)
 645 only needs to store information about which wallet subkeys have been active. In terms of
 646 communication, our protocol transmits only a constant number of elements when using a
 647 succinct proof system. This compares favorably to solutions such as SPV which incur linear
 648 costs or even solutions based on succinct proof systems for the chain tip, but where the LC
 649 needs to verify transaction inclusion into blocks via Merkle proofs.

650 **Transaction discovery via HD Wallets.** Assuming that the client only uses a single
 651 address can be unrealistic, as is requiring the LC to send the SP a list of addresses presumed
 652 active. Our protocol is compatible with BIP-32 [49] hierarchical deterministic (HD) wallets.
 653 This enables child addresses to be deterministically created from a parent address by hashing
 654 the parent key X , a chaincode value cc and an address index i . Concretely, $o_i := \mathsf{H}(X, cc, i)$
 655 is the private key offset of child key i , i.e. the i -th child keypair is $x_i = x + o_i$ and $X_i = X \cdot g^{o_i}$.
 656 Child chaincodes are determined in the same way.

657 Thus, an LC can simply transmit the chaincode corresponding to its public key X , and the
 658 SP will be able to derive a list of child addresses (bounding indexes can be done heuristically,
 659 or with hints from the client). The protocol then proceeds with minor changes:

- 660 1. The challenge c is now derived as $c = \mathsf{H}(R, X_c, m)$
 - 661 2. The witness w now includes the index of the child key i (or vector of indices \mathbf{i} and chaincodes
 662 \mathbf{cc} if the child is further down the tree).
 - 663 3. The relation now checks that the hashed value is of the form $X_c = X \cdot g^{o_c}$ where $o_c =$
 664 $\mathsf{HDDerive}(X, \mathbf{i}, \mathbf{cc})$.
 - 665 4. The SP modifies the component $s := xc + r$ to $s' = s + o_c \cdot c$, adjusting for the child offset.
- 666 Notably, the LC does not learn which child address was used, it can simply sign with regards
 667 to the sent public key X and the SP can maul the signature as needed.
- 668 **Asset and Non-Asset Compensation.** In our model, we have so far assumed that an
 669 SP performs its services in exchange for a fee. This fee may be either flat, or calculated on

670 the basis of transaction size, or complexity of transaction specification, etc.—the details of
 671 this fee calculation are dependent on the specifics of the implementation of our protocol, the
 672 marketplace, and SPs preferences. The fee may also be specified in either an amount of primary
 673 asset tokens, or some other user-defined tokens. A fee requirement in user-defined tokens could
 674 be useful in the case of the SP service for LCs being associated with another type of service
 675 trading in such tokens, such as a videogame token marketplace. Yet another option for SPs
 676 is to request non-monetary compensation for their services, such as requiring engagement with
 677 a specific smart contract (i.e. SP does DApp fee sponsorship), voting for a specific update,
 678 delegating stake, etc. All these actions can be performed by the very transaction that SP
 679 constructed according to LC’s specification.

680 9 Implementation

681 To assess the efficiency of our construction, we implement and benchmark the **NArg** component
 682 of the Cavefish protocol. Given today’s implementations of zk-SNARKs, running the **NArg**
 683 is expected to be the most time and resource intensive part of our light client protocol.

684 We base our tests on the trusted-setup zk-SNARK system *Groth16* [27] implemented by
 685 *Iden3* [30]. The circuits we construct and benchmark are inspired by the implementation of [24]
 686 and written in the domain-specific language of Circom 2.1.

687 The **NArg** we implement captures the relation R_{Cavefish} in the protocol in Section 7.2. We
 688 implement the “light” version of Cavefish, i.e., the intent is given by
 689 $\text{int}_{\text{light}}(\text{tx}) := \{\text{return } (\text{mkAbs}(\text{tx}) = \text{mkAbs}(\text{tx}_0))\}$. The light client effectively specifies the
 690 transaction, except for the input UTxOs and note field **aux**.

691 9.1 Implementation choices

692 To implement and measure the arithmetic complexity of the relation R_{Cavefish} , we use BN254
 693 as the curve for *Groth16* to operate on, i.e., the curve group has 254 bits and the relation is
 694 instantiated over an arithmetic circuit with modulus of 254 bits. The BN254 curve is one of
 695 the standard choices in Circom and forces the inputs to the circuit to be elements of the field
 696 given by BN254. We capture transaction tx and abstract transaction tx_A as bit strings of some
 697 length n encoded as field elements, which allows us to instantiate a circuit that can handle
 698 transactions of length up to n .

699 In order to implement the encryption scheme PKE that is needed to encrypt tx as C , we use
 700 ElGamal [21] public key encryption over the Baby-JubJub curve [48]. The reason for choosing
 701 Baby-JubJub curve is that the field given by BN254 is the base field of the Baby-JubJub curve
 702 and thus, any element can be represented as two elements of the BN254 field. As a consequence,
 703 the group operation is efficiently arithmetizable in the circuit [24]. Furthermore, we encapsulate
 704 key and encryption of ElGamal using DHIES [1], i.e., the shared secret in ElGamal serves as a
 705 seed to the PRF generating random group elements for an additive one-time-pad that encrypts
 706 tx . The PRF is instantiated as a Poseidon hash [26] that is efficiently arithmetizable by design.

707 Unfortunately, most common cryptographic hash functions are not “circuit-friendly” [26]
 708 and thus it would be beneficial to optimize the hash function used to create the Schnorr
 709 challenge in the blind signature protocol. However, to be compatible with existing blockchain
 710 platforms and create standard Schnorr signatures, we have to adopt the exact hash function
 711 specified by the respective ledger. We use the library of hash functions in [33] which provides
 712 Circom implementations for many popular hash functions. The authors put effort into the
 713 optimization, but more efficient implementations might be possible. Despite potential further

714 improvements, the complexity of our resulting circuit is mainly governed by the number of
715 rounds the hash function has to execute when producing the Schnorr challenge.

716 9.2 Benchmarks

717 We test and benchmark partially blind signatures for both Bitcoin and Cardano. Partial
718 blindness can be interpreted as a predicate itself (see Sections 6.3 and 7.2). The “light” version
719 of Cavefish uses int_{light} which can be implemented with substring equality checks, provided the
720 transaction to be signed can be unambiguously separated into blinded and nonblinded parts.
721 Cardano transactions are CBOR-encoded and contain the transaction-id and index for every
722 UTxO serving as input [31]. Similarly, Bitcoin transactions (after Taproot update [19]) feature
723 an input count followed by a list of inputs in binary format, each consisting of a transaction
724 hash and output index. We assume for our benchmarks that one UTxO serves as input to the
725 transaction and needs to be blinded.

726 We measure the number of constraints, the proving key size, and the proof size that $R_{Cavefish}$
727 requires in Circom. The number of constraints are obtained from the arithmetization when
728 given as a R1CS relation. We also keep track of the time it takes to create the resulting circuit,
729 the proving time and the proof verification time.

730 The results are summarized in Table 1 for Bitcoin and Cardano as the target platform. The
731 experiments were executed on commodity hardware based on an Intel(R) Core(TM) i7-8750H
732 CPU operating at 2.20 GHz with 12 cores and 16 GB of RAM.

 **Table 1** Benchmark of $R_{Cavefish}$ for Bitcoin and Cardano implementing the “light” version of Cavefish.

	Bitcoin	Cardano
Signature scheme	Schnorr	EdDSA
Curve	Secp256k1	Ed25519
Hash	SHA-256	SHA-512
Transaction size	254 B 288 b blinded	285 B 333 b blinded
Proving key size	115 MB	116 MB
Proving key verification time	18.6 s	20.5 s
Verification key size	67 kB	93 kB
Proving time	5.1 s	5.8 s
Proof size	806 B	805 B
Proof verification time	0.57 s	0.59 s
Number of constraints	226509	245181

733 Bitcoin has one major conceptual difference from Cardano as the message being signed by
734 the Schnorr signature scheme is the hash of the transaction $H(tx)$ instead of the transaction
735 itself. Therefore, $WBPS[P]$ is executed with $m := H(tx||aux_{nt})$. A straightforward adaption to
736 $WBPS[P]$ that supports hashed transactions is to allow the predicate P to accept an additional
737 input which is a witness attesting to the signed message m satisfying P . This additional witness
738 must be included in the witness ω for $R_{Cavefish}$ [24]. Therefore, compared to Cardano, the
739 circuit for Bitcoin has to perform an additional invocation of the hash function (SHA-256).
740 On the other hand, the hash function used by Cardano (SHA-512) has a larger codomain and
741 higher complexity, which requires more constraints.

742 We remark that the projects our implementation builds upon are in development, such
743 as the hash function library, and lack certain optimizations. Nevertheless, the introduction
744 of a weak variant of blindness in our light client protocol allowed us to obtain results within
745 practical bounds and shows that the Cavefish protocol is deployable in the real world.

746 **References**

- 747 1 Michel Abdalla, Mihir Bellare, and Phillip Rogaway. Dhies: An encryption scheme based on
748 the diffie-hellman problem. *IACR Cryptol. ePrint Arch.*, 1999:7, 1999.
- 749 2 Frederik Armknecht, Ghassan Karamé, Malcom Mohamed, and Christiane Weis. Practical light
750 clients for committee-based blockchains, 2024. URL: <https://arxiv.org/abs/2410.03347>,
751 arXiv:2410.03347.
- 752 3 Foteini Baldimtsi, Varun Madathil, Alessandra Scafuro, and Linfeng Zhou. Anonymous
753 Lottery In The Proof-of-Stake Setting . In *2020 IEEE 33rd Computer Security Foundations
754 Symposium (CSF)*, pages 318–333, Los Alamitos, CA, USA, June 2020. IEEE Computer
755 Society. URL: <https://doi.ieee.org/10.1109/CSF49147.2020.00030>,
756 doi:10.1109/CSF49147.2020.00030.
- 757 4 Bellare, Nampempre, Pointcheval, and Semanko. The one-more-rsa-inversion problems and
758 the security of chaum’s blind signature scheme. *Journal of Cryptology*, 16:185–215, 2003.
- 759 5 Fabrice Benhamouda, Tancrede Lepoint, Julian Loss, Michele Orrù, and Mariana Raykova.
760 On the (in)security of ROS. *Cryptology ePrint Archive*, Paper 2020/945, 2020. URL:
761 <https://eprint.iacr.org/2020/945>, doi:10.1007/s00145-022-09436-0.
- 762 6 Daniel J. Bernstein, Niels Duif, Tanja Lange, Peter Schwabe, and Bo-Yin Yang. High-
763 speed high-security signatures. *Cryptology ePrint Archive*, Paper 2011/368, 2011. URL:
764 <https://eprint.iacr.org/2011/368>.
- 765 7 Jan Bobolz, Jesus Diaz, and Markulf Kohlweiss. Foundations of anonymous signatures: Formal
766 definitions, simplified requirements, and a construction based on general assumptions. In
767 Jeremy Clark and Elaine Shi, editors, *Financial Cryptography and Data Security*, pages 121–139,
768 Cham, 2025. Springer Nature Switzerland.
- 769 8 Dan Boneh and Chelsea Komlo. Threshold signatures with private accountability. *Cryptology
770 ePrint Archive*, Paper 2022/1636, 2022. URL: <https://eprint.iacr.org/2022/1636>,
771 doi:10.1007/978-3-031-15985-5_19.
- 772 9 Benedikt Bünz, Lucianna Kiffer, Loi Luu, and Mahdi Zamani. Flyclient: Super-light
773 clients for cryptocurrencies. *Cryptology ePrint Archive*, Paper 2019/226, 2019. URL:
774 <https://eprint.iacr.org/2019/226>.
- 775 10 Cardano Team. Full Cardano Ledger. [https://intersectmbo.github.io/
776 formal-ledger-specifications/pdfs/cardano-ledger.pdf](https://intersectmbo.github.io/formal-ledger-specifications/pdfs/cardano-ledger.pdf), 2024.
- 777 11 Pyrrros Chaidos and Aggelos Kiayias. Mithril: Stake-based threshold multisignatures. In
778 Markulf Kohlweiss, Roberto Di Pietro, and Alastair Beresford, editors, *Cryptology and Network
779 Security*, pages 239–263, Singapore, 2025. Springer Nature Singapore.
- 780 12 Rutchathon Chairattana-Apirom, Lucjan Hanzlik, Julian Loss, Anna Lysyanskaya, and Benedikt
781 Wagner. Pi-cut-choo and friends: Compact blind signatures via parallel instance cut-and-choose
782 and more. In *Annual International Cryptology Conference*, pages 3–31. Springer, 2022.
- 783 13 Manuel M. T. Chakravarty, James Chapman, Kenneth MacKenzie, Orestis Melkonian, Jann
784 Müller, Michael Peyton Jones, Polina Vinogradova, Philip Wadler, and Joachim Zahnentferner.
785 UTXO_{ma}: UTXO with multi-asset support. In Tiziana Margaria and Bernhard Steffen, editors,
786 *Leveraging Applications of Formal Methods, Verification and Validation: Applications - 9th
787 International Symposium on Leveraging Applications of Formal Methods, ISoLA 2020, Rhodes,
788 Greece, October 20-30, 2020, Proceedings, Part III*, volume 12478 of *Lecture Notes in Computer
789 Science*, pages 112–130. Springer, 2020. doi:10.1007/978-3-030-61467-6_9.
- 790 14 Manuel MT Chakravarty, Sandro Coretti, Matthias Fitzi, Peter Gaži, Philipp Kant, Aggelos
791 Kiayias, and Alexander Russell. Fast isomorphic state channels. In *Financial Cryptography
792 and Data Security: 25th International Conference, FC 2021, Virtual Event, March 1–5, 2021,
793 Revised Selected Papers, Part II 25*, pages 339–358. Springer, 2021.
- 794 15 Panagiotis Chatzigiannis, Foteini Baldimtsi, and Konstantinos Chalkias. SoK:
795 Blockchain light clients. *Cryptology ePrint Archive*, Paper 2021/1657, 2021. URL:
796 <https://eprint.iacr.org/2021/1657>.

- 797 **16** David Chaum and Torben Pryds Pedersen. Wallet databases with observers. In Ernest F.
 798 Brickell, editor, *Advances in Cryptology — CRYPTO' 92*, pages 89–105, Berlin, Heidelberg,
 799 1993. Springer Berlin Heidelberg.
- 800 **17** Adrian Brink Christopher Goes, Awa Sun Yin. Anoma : a unified architecture for full-stack
 801 decentralized applications, 2022. URL: <https://github.com/anoma/whitepaper/blob/main/whitepaper.pdf>.
- 803 **18** Arthur Britto David Schwartz, Noah Youngs. The ripple protocol consensus algorithm, 2018.
 804 URL: https://ripple.com/files/ripple_consensus_whitepaper.pdf.
- 805 **19** Bitcoin Developers. Bip 340: Schnorr signatures for secp256k1. <https://github.com/bitcoin/bips/blob/master/bip-0340.mediawiki>, 2021. Accessed: 2025-05-28.
- 807 **20** Ankur Dubey. A decentralised solver architecture for executing intents on evm blockchain. <https://ethresear.ch/t/a-decentralised-solver-architecture-for-executing-intents-on-evm-blockchain/16608>, 2023. Accessed: 2025-05-28.
- 811 **21** Taher ElGamal. A public key cryptosystem and a signature scheme based on discrete logarithms.
IEEE transactions on information theory, 31(4):469–472, 1985.
- 813 **22** Ergo Team. Ergo: A Resilient Platform For Contractual Money. <https://whitepaper.io/document/753/ergo-1-whitepaper>, 2019.
- 815 **23** Georg Fuchsbauer, Antoine Plouviez, and Yannick Seurin. Blind schnorr signatures and signed
 816 elgamal encryption in the algebraic group model. In *Advances in Cryptology – EUROCRYPT 2020: 39th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Zagreb, Croatia, May 10–14, 2020, Proceedings, Part II*, page 63–95, Berlin,
 817 Heidelberg, 2020. Springer-Verlag. doi:10.1007/978-3-030-45724-2_3.
- 820 **24** Georg Fuchsbauer and Mathias Wolf. Concurrently secure blind schnorr signatures. In Marc
 821 Joye and Gregor Leander, editors, *Advances in Cryptology – EUROCRYPT 2024*, pages 124–160,
 822 Cham, 2024. Springer Nature Switzerland.
- 823 **25** Ariel Gabizon, Zachary J Williamson, and Oana Ciobotaru. Plonk: Permutations over lagrange-
 824 bases for oecumenical noninteractive arguments of knowledge. *Cryptology ePrint Archive*, 2019.
- 825 **26** Lorenzo Grassi, Dmitry Khovratovich, Christian Rechberger, Arnab Roy, and Markus Schoenberger.
 826 Poseidon: A new hash function for Zero-Knowledge proof systems. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 519–535. USENIX Association, August 2021.
 827 URL: <https://www.usenix.org/conference/usenixsecurity21/presentation/grassi>.
- 829 **27** Jens Groth. On the size of pairing-based non-interactive arguments. *Cryptology ePrint Archive*,
 830 Paper 2016/260, 2016. URL: <https://eprint.iacr.org/2016/260>.
- 831 **28** Lucjan Hanzlik, Julian Loss, and Benedikt Wagner. Rai-choo! evolving blind signatures to the
 832 next level. In *Annual International Conference on the Theory and Applications of Cryptographic
 833 Techniques*, pages 753–783. Springer, 2023.
- 834 **29** Dominik Hartmann and Eike Kiltz. Limits in the provable security of ECDSA signatures.
 835 Cryptology ePrint Archive, Paper 2023/914, 2023. URL: <https://eprint.iacr.org/2023/914>.
- 836 **30** iden3. Circom: Circuit compiler for zero-knowledge proofs. <https://github.com/iden3/circom>,
 837 2023. Accessed: 2025-05-28.
- 838 **31** IntersectMBO. Cardano ledger specifications and implementations. <https://github.com/IntersectMBO/cardano-ledger>, 2025. Accessed: 2025-05-28.
- 840 **32** Joseph Poon and Thaddeus Dryja. The Bitcoin Lightning Network: Scalable Off-Chain Instant
 841 Payments. Technical report, University of Zurich, Department of Informatics, 01 2010.
- 842 **33** Balazs Komuves. hash-circuits: Hashing circuits implemented in circom. <https://github.com/bkomuves/hash-circuits>, 2025. MIT License, maintained by Faulhorn Labs.
- 844 **34** Yuan Lu, Qiang Tang, and Guiling Wang. Generic superlight client for permissionless blockchains.
 845 Cryptology ePrint Archive, Paper 2020/844, 2020. URL: <https://eprint.iacr.org/2020/844>.
- 846 **35** Giulio Malavolta, Pedro Moreno-Sanchez, Aniket Kate, Matteo Maffei, and Srivatsan Ravi.
 847 Concurrency and privacy with payment-channel networks. In *Proceedings of the 2017 ACM*

- 848 *SIGSAC Conference on Computer and Communications Security, CCS '17*, pages 455–471, New
 849 York, NY, USA, 2017. Association for Computing Machinery. doi:10.1145/3133956.3134096.
- 850 36 Gregory Maxwell, Andrew Poelstra, Yannick Seurin, and Pieter Wuille. Simple schnorr
 851 multi-signatures with applications to bitcoin. *Des. Codes Cryptography*, 87(9):2139–2164,
 852 September 2019. doi:10.1007/s10623-019-00608-x.
- 853 37 S. Nakamoto. Bitcoin: A Peer-to-Peer Electronic Cash System. <https://bitcoin.org/en/bitcoin-paper>, October 2008.
- 854 38 David Pointcheval and Jacques Stern. Security arguments for digital signatures and blind
 855 signatures. *J. Cryptology*, 13:361–396, 2000. doi:10.1007/s001450010003.
- 856 39 C. P. Schnorr. Efficient identification and signatures for smart cards. In Gilles Brassard, editor,
 857 *Advances in Cryptology — CRYPTO' 89 Proceedings*, pages 239–252, New York, NY, 1990.
 858 Springer New York.
- 859 40 Claus Peter Schnorr. Security of blind discrete log signatures against interactive attacks. In
 860 Sihan Qing, Tatsuaki Okamoto, and Jianying Zhou, editors, *Information and Communications
 861 Security*, pages 1–12, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 862 41 Ertem Nusret Tas, David Tse, Lei Yang, and Dionysis Zindros. Light clients for lazy blockchains,
 863 2024. URL: <https://arxiv.org/abs/2203.15968>, arXiv:2203.15968.
- 864 42 Ethereum Team. Ethereum development documentation, 2023. URL: <https://ethereum.org/en/developers/docs/>.
- 865 43 Khalani Team. Khalani arcadia: An open platform for intent-driven agent collaboration, 2025.
 866 URL: https://khalani.network/aip_whitepaper.pdf.
- 867 44 Raiden Team. Raiden network 3.0.1 documentation, 2025. URL: <https://raiden-network.readthedocs.io/>.
- 868 45 Stefano Tessaro and Chenzhi Zhu. Short pairing-free blind signatures with exponential
 869 security. In *Annual International Conference on the Theory and Applications of Cryptographic
 870 Techniques*, pages 782–811. Springer, 2022.
- 871 46 Psi Vesely, Kobi Gurkan, Michael Straka, Ariel Gabizon, Philipp Jovanovic, Georgios
 872 Konstantopoulos, Asa Oines, Marek Olszewski, and Eran Tromer. Plumo: An ultralight
 873 blockchain client, 2025. White paper. URL: <https://celo.org/papers/plumo>.
- 874 47 David Wagner. A generalized birthday problem. In Moti Yung, editor, *Advances in Cryptology
 875 — CRYPTO 2002*, pages 288–304, Berlin, Heidelberg, 2002. Springer Berlin Heidelberg.
- 876 48 Barry WhiteHat, Jordi Baylina, and Marta Bellés. Baby jubjub elliptic curve. *Ethereum
 877 Improvement Proposal, EIP-2494*, 29, 2020.
- 878 49 Pieter Wuille. BIP 0032: Hierarchical deterministic wallets. https://github.com/bitcoin/bips/blob/master/bip-0032.mediawiki#user-content-Full_wallet_sharing_m, 2012.
- 879 50 ScaleSphere Foundation Ltd. (“Foundation”). Celer network: Bring internet scale to every
 880 blockchain, 2018. URL: <https://celer.network/doc/CelerNetwork-Whitepaper.pdf>.

885 **A Expanded Technical Backgroud**

886 **A.1 Discrete Log Groups**

887 ► **Definition 11.** A group generator GGen is a probabilistic polynomial time (p.p.t.) algorithm
 888 with input a security parameter λ and outputs a group description \mathbb{G}, g, q such that \mathbb{G} is a group
 889 of prime order $q \approx 2^\lambda$ with generator g . We say that the discrete logarithm problem is hard w.r.t.
 890 GGen if for all p.p.t \mathcal{A} we have that

$$\Pr[(\mathbb{G}, g, q) \leftarrow \text{GGen}(1^\lambda); t \leftarrow \mathbb{Z}_q; h \leftarrow g^t : t = A(\mathbb{G}, g, q, h)] \text{ is negligible in } \lambda.$$

891 **A.2 Public Key encryption and Schnorr Signatures**

- 892 A public key encryption scheme PKE comprises a set of polynomial time algorithms $\text{KeyGen}, \text{Enc}, \text{Dec}$
 893 with the following syntax:
- 894 ■ $\text{KeyGen}(1^\lambda) \rightarrow (ek, dk)$. Creates a public/private keypair ek, dk . The encryption key ek also
 895 defines the message space \mathcal{M}
 - 896 ■ $\text{Enc}(ek, m; \rho) \rightarrow C$. Encrypts a message $m \in \mathcal{M}$ under the public encryption key ek .
 - 897 ■ $\text{Dec}(dk, C) \rightarrow m$. Decrypts a ciphertext C using the private decryption key dk .

► **Definition 12.** A public key encryption scheme PKE is correct if

$$\Pr[(ek, dk) \leftarrow \text{PKE.KeyGen}(1^\lambda); m \leftarrow \mathcal{M} : \text{PKE.Dec}(sk, \text{PKE.Enc}(pk, m)) = m] = 1.$$

► **Definition 13.** A public key encryption scheme PKE is IND-CPA secure if for all stateful
 p.p.t. adversaries \mathcal{A} , the difference

$$\left| \Pr[(ek, dk) \leftarrow \text{PKE.KeyGen}(1^\lambda); (m_0, m_1) \leftarrow A(ek); d \leftarrow \{0, 1\}; c^* \leftarrow \text{PKE.Enc}(ek, m_d) : A(c^*) = d] - \frac{1}{2} \right|$$

898 is negligible in λ .

899 ► **Definition 14.** The Schnorr digital signature scheme SDS is defined for a group generator
 900 GGen and a hash function generator HGen and operates as shown on Figure 7.

SDS.Setup(1^λ)	SDS.KeyGen(sp)
$(\mathbb{G}, g, q) \leftarrow \text{GGen}(1^\lambda)$	$(\mathbb{G}, g, q) : \subseteq sp$
$H \leftarrow \text{HGen}(q)$	$x \leftarrow \mathbb{Z}_q; X \leftarrow g^x$
$sp \leftarrow (\mathbb{G}, g, q, H)$	$sk, vk \leftarrow (par, x), (par, X)$
return sp	return (sk, vk)
SDS.Sign(sk, m)	SDS.Ver(vk, m, σ)
$(\mathbb{G}, g, q, H, x) : \subseteq sk$	$(\mathbb{G}, g, q, H, X) : \subseteq vk$
$r \leftarrow \mathbb{Z}_q; R \leftarrow g^x; c \leftarrow H(R, X, m)$	$(R, s) \leftarrow \sigma; c \leftarrow H(R, X, m)$
$s \leftarrow (r + c \cdot x) \bmod q$	return $(g^s = R \cdot X^c)$
return $\sigma \leftarrow (R, s)$	

► **Figure 7** The Schnorr signature scheme SDS, with group and hash generators GGen, HGen.

► **Definition 15.** A digital signature scheme DS is correct if

$$\Pr[sp \leftarrow \text{DS.Setup}(1^\lambda); (sk, vk) \leftarrow \text{DS.KeyGen}(sp); m \leftarrow \mathcal{M}_S : \text{DS.Ver}(vk, m, \text{Sign}(sk, m)) = 1] = 1.$$

► **Definition 16.** A digital signature scheme DS is strongly existentially unforgeable against chosen message attacks (sEUF-CMA) if for all p.p.t. adversaries \mathcal{A} , we have

$$\Pr[s\text{EUF-CMA}_{\text{DS}}^{\mathcal{A}}(1^\lambda)]$$

901 is negligible in λ where the game $G_{\text{sEUF-CMA}}^{\mathcal{A}}_{\text{DS}}$ is defined in Figure 8.

GsEUF-CMA_{DS}^A(1^λ) $sp \leftarrow DS.\text{Setup}(1^\lambda); Q \leftarrow \emptyset$ $(sk, vk) \leftarrow DS.\text{KeyGen}(sp)$ $(m^*, \sigma^*) \leftarrow \mathcal{A}^{\text{OSig}}(vk)$ return $(m^*, \sigma^*) \neq Q \wedge DS.\text{Ver}(vk, m^*, \sigma^*)$	OSig(m) $\sigma \leftarrow DS.\text{Sign}(skk, m)$ $Q \leftarrow Q \cup \{(m, s)\}$ return σ
--	--

■ **Figure 8** The security experiment $s\text{EUF-CMA}_{DS}^A$ and supporting oracle OSig .

902 A.3 Parametrized Non-Interactive Zero Knowledge Arguments

903 Following [24], we define non-interactive zero knowledge arguments with regards to parametrized
904 polynomial relations $\mathcal{P}: \{0,1\}^* \times \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$, where the first argument represents
905 a parameter set par (e.g. a group description). Given a value of par , we say that w is a witness
906 for statement θ if $\mathcal{P}(par, \theta, w) = 1$, i.e. $R = R_{par}(\theta, w) := \mathcal{P}(par, \theta, w)$ is an NP-relation, and
907 $\mathcal{L} = \mathcal{L}_{par}$ is an NP-language. A NIZK for a relation \mathcal{P} operates as follows:

- 908 ■ $\text{Rel}(1^\lambda) \rightarrow par$. Generates a parameter set par that defines R and \mathcal{L} .
- 909 ■ $\text{Setup}(par) \rightarrow (\text{crs}, \tau)$. Generates a common reference string (CRS) and trapdoor τ used
910 by the simulator. We assume the CRS contains a description of R .
- 911 ■ $\text{Prove}(\text{crs}, \theta, w) \rightarrow \pi$. Given a CRS crs , statement θ and witness w for θ , produces a proof π .
- 912 ■ $\text{Ver}(\text{crs}, \theta, \pi) \rightarrow \{0,1\}$. Given a CRS crs , a statement θ and a proof π outputs 1 or 0, accepting
913 or rejecting the proof.
- 914 ■ $\text{SimProve}(\text{crs}, \theta, \tau)$. Given a CRS crs , statement θ and trapdoor τ for crs , produces a
915 simulated proof π .

916 ► **Definition 17.** A system $\text{NArg}[R]$ is perfectly correct if for all unbounded adversaries \mathcal{A}

$$917 \quad \Pr[par \leftarrow \text{NArg}.\text{Rel}(1^\lambda); (crs, \tau) \leftarrow \text{NArg}.\text{Setup}(par); (\theta, w) \leftarrow A(crs): \\ 918 \quad \neg R(\theta, w) \vee \text{NArg}.\text{Ver}(crs, \theta, \text{NArg}.\text{Prove}(crs, \theta, w))] = 1$$

919

920 ► **Definition 18.** A system $\text{NArg}[R]$ is adaptably computationally sound if for all p.p.t. ad-
921 versaries \mathcal{A}

$$922 \quad \Pr[par \leftarrow \text{NArg}.\text{Rel}(1^\lambda); (crs, \tau) \leftarrow \text{NArg}.\text{Setup}(par); (\theta, \pi) \leftarrow A(crs): \\ 923 \quad \neg L(\theta) \wedge \text{NArg}.\text{Ver}(crs, \theta, \pi)] \text{ is negligible in } \lambda.$$

924 ► **Definition 19.** A system $\text{NArg}[R]$ is computationally zero-knowledge if for all p.p.t. adver-
925 saries \mathcal{A}

$$926 \quad |\Pr[par \leftarrow \text{NArg}.\text{Rel}(1^\lambda); (crs, \tau) \leftarrow \text{NArg}.\text{Setup}(par); d \leftarrow \{0,1\}: \\ 927 \quad d = A^{\text{OProve}_d}(crs)] - \frac{1}{2}| \text{ is negligible in } \lambda, \text{ where} \\ 928 \quad \text{OProve}_0(\theta, w) := \text{if } \neg R(\theta, w) \text{ return } \perp; \text{ return } \text{NArg}.\text{Prove}(\theta, w), \text{ and} \\ 929 \quad \text{OProve}_1(\theta, w) := \text{if } \neg R(\theta, w) \text{ return } \perp; \text{ return } \text{NArg}.\text{SimProve}(\theta, \tau).$$

930 **B Additional Ledger Syntax**

931 In Figure 9 we introduce the standard ledger syntax that we use throughout.

$\mathbb{H} = \bigcup_{n=0}^{\infty} \{0,1\}^{8n}$	the type of bytestrings
$(a,b) : \text{Interval}[A]$	intervals over a totally-ordered set A
$\text{Key} \mapsto \text{Value} \subseteq \{ k \mapsto v \mid k \in \text{Key}, v \in \text{Value} \}$	finite map with unique keys
$[a_1; \dots; a_k] : [C]$	finite list with terms of type C
$h :: t : [C]$	list with head h and tail t
$x \cup \text{nothing} : A^?$	maybe type over A
$a \{ \text{field} = x \} : A$	record of type A with field changed to x

932 **Figure 9** Notation

933 Figure 10 lists the primitives and derived types that comprise the foundations of the EUTxO
 934 model, along with some ancillary definitions. (Outputs normally refer to transaction IDs by
 hash, but we simplify here for clarity.)

LEDGER PRIMITIVES

$\text{checkSig} : \text{Tx} \rightarrow \text{PubKey} \rightarrow \mathbb{H} \rightarrow \{0,1\}$
checks that a given key signed a transaction

HELPER FUNCTIONS

$\text{txid} : \text{Tx} \rightarrow \text{TxId}$
 $\text{txid } tx = \text{hash} (tx \{ \text{sigKeys} = \text{dom} (tx.\text{sigs}), \text{sigs} = \emptyset \})$

$\text{toMap} : \mathbb{N} \rightarrow [\text{TxOut}] \rightarrow (\mathbb{N} \mapsto \text{TxOut})$
 $\text{toMap}(_, []) = []$
 $\text{toMap}(ix, u :: outs) = \{ ix \mapsto u \} \cup \text{toMap}(ix+1, outs)$
constructs a map from a list of outputs

$\text{mkOuts} : \text{Tx} \rightarrow \text{UTxO}$
 $\text{mkOuts}(tx) = \{ (tx, ix) \mapsto o \mid (ix \mapsto o) \in \text{toMap}(0, tx.\text{outputs}) \}$
constructs a UTxO set from a list of outputs of a given transaction

$\text{MOF} : \mathbb{N} \rightarrow \mathbb{N} \rightarrow (A \rightarrow \{0,1\}) \rightarrow [A] \rightarrow \{0,1\}$
 $\text{MOF } k m f [] = m \leq k$
 $\text{MOF } k m f (h :: t) = \text{if } (m \leq k) \text{ then } 1 \text{ else } (\text{MOF } (k + a) m f t)$
where $a = \text{if } (f(h)) \text{ then } 1 \text{ else } 0$
returns 1 if enough elements of a list satisfy given function

■ **Figure 10** Primitives and basic types for the UTxO_{ma} model

935 **B.1 Transaction Validation Rules**

936 A transaction tx is *valid* if it follows the following rules.

937 (i) **The transaction has at least one input:**

938 $tx.\text{inputs} \neq \{ \}$

939 (ii) **The current slot is within transaction validity interval:**

940 $slot \in tx.\text{validityInterval}$

941 (iii) **All outputs have positive values:**

942 $\forall o \in tx.\text{outputs}, o.\text{value} > \emptyset$

943 (iv) **All output references of transaction inputs exist in the UTxO:**

944 $tx.\text{inputs} \subseteq \text{dom utxo}$

945 (v) **Value is preserved:**

946 $tx.\text{mint} + \sum_{i \in tx.\text{inputs}, (i \mapsto o) \in utxo} o.\text{value} = \sum_{o \in tx.\text{outputs}} o.\text{value} + \text{toValue}(tx.\text{fee})$

CONSTRUCTORS OF Script

RequireMOF	$: \mathbb{N} \rightarrow [\text{Script}]$	$\rightarrow \text{Script}$
RequireSig	$: \text{PubKey}$	$\rightarrow \text{Script}$
RequireTimeStart	$: \text{Slot}$	$\rightarrow \text{Script}$
RequireTimeExpire	$: \text{Slot}$	$\rightarrow \text{Script}$

EVALUATION OF Script

$$\begin{aligned} \llbracket _ \rrbracket & : \text{Script} \rightarrow ((\text{SetPubKey}) \times (\text{Slot} \times \text{Slot})) \rightarrow \{0,1\} \\ \llbracket \text{RequireMOF } n \text{ ls} \rrbracket(khs, (t1, t2)) & = \text{MOF } 0 \text{ n } (\llbracket _ \rrbracket(khs, (t1, t2))) \text{ ls} \\ \llbracket \text{RequireSig } k \rrbracket(khs, (t1, t2)) & = k \in khs \\ \llbracket \text{RequireTimeStart } t1' \rrbracket(khs, (t1, t2)) & = t1' \leq t1 \\ \llbracket \text{RequireTimeExpire } t2' \rrbracket(khs, (t1, t2)) & = t2 \leq t2' \end{aligned}$$

Figure 11 Script constructors and evaluation

947 (vii) All inputs validate:

948 $\forall i \in tx.\text{inputs}, i \mapsto (s, v) \in utxo, \llbracket s \rrbracket(\text{dom } (tx.\text{sig}), tx.\text{validityInterval}) = 1$

949 (ix) All minting scripts validate:

950 $\forall p \mapsto _ \in tx.\text{mint}, \llbracket p \rrbracket(\text{dom } (tx.\text{sig}), tx.\text{validityInterval}) = 1$

951 (x) All signatures are correct:

952 $\forall (pk \mapsto s) \in tx.\text{sig}, \text{checkSig}(tx, pk, s) = 1$

953 (i) The fee is sufficient:

954 $s.\text{minfee} \leq tx.\text{fee}$

955 **B.2 Script construction and Evaluation**956 In Figure 11 we present the constructors and evaluation rules for scripts, and in Figure 12 we
957 explain the transaction building function `mkToSpec`.

TRIVIAL TRANSACTION $\text{initTx}_{a,mf}$

```
 $\text{initTx}_{a,mf} = \{ \text{inputs} = \emptyset,$ 
 $\text{outputs} = [(a, \text{tip})],$ 
 $\text{validityInterval} = [\text{nothing}, \text{nothing}],$ 
 $\text{mint} = 0,$ 
 $\text{fee} = mf$ 
 $\text{aux} = []$ 
 $\text{sigs} = \emptyset \}$ 
```

INPUT SELECTION FUNCTION

```
 $\text{mkIns} : (\text{LState} \times (\text{SetTxIn}) \times \text{Value} \times \text{Script}) \rightarrow (\text{SetTxIn})$ 
 $\text{mkIns}(l, i, v, s) = \text{if } \neg(v \leq 0) \text{ then}$ 
 $\quad \text{if } (v > 0) \text{ then}$ 
 $\quad \quad \text{mkIns}(l \setminus (\{j \mapsto \_\}), i \cup j, v - (u(j).value), s)$ 
 $\quad \text{else } i$ 
 $\text{else}$ 
 $\quad \text{nothing}$ 
 $\text{where}$ 
 $\quad j = \text{pickInput } l \ s$ 
 $\quad u = \text{utxo } l$ 
```

AUXILIARY $\text{mkToSpec}'$ DEFINITION

```
 $\text{mkToSpec}' : \text{LState} \rightarrow \mathcal{I}_{\text{post}} \rightarrow \text{Tx} \rightarrow \text{Tx}^?$ 
 $\text{mkToSpec}' l (\text{MustMint } v)(tx) = tx \{ \text{mint} = v + tx.\text{mint}, \text{sigs} = tx.\text{sigs} \cup \text{getSigsVal } v,$ 
 $\quad \text{validityInterval} = \text{restrictIntervalVal } tx.\text{validityInterval } v \}$ 
 $\text{mkToSpec}' l (\text{SpendFrom } s)(tx) = tx \{ \text{outputs} = tx.\text{inputs}$ 
 $\quad \cup \text{newIns}, \text{sigs} = tx.\text{sigs} \cup \text{getSigsUTxO newIns } l,$ 
 $\quad \text{validityInterval} =$ 
 $\quad \quad \text{restrictIntervalUTxO } tx.\text{validityInterval newIns } l \}$ 
 $\quad \text{where}$ 
 $\quad \text{newIns} = \text{mkIns } l (tx.\text{inputs}) \text{ (produced - consumed)} \ s$ 
 $\text{mkToSpec}' l (\text{MaxInterval } i)(tx) = tx \{ \text{validityInterval} = (l.\text{slot},$ 
 $\quad \min \{l.\text{slot} + i, tx.\text{validityInterval}_2\}) \}$ 
 $\text{mkToSpec}' l (\text{PayTo } (s, v))(tx) = tx \{ \text{outputs} = tx.\text{outputs} \cup (s, v) \}$ 
 $\text{mkToSpec}' l (\text{ChangeTo } s)(tx) = \text{if consumed - produced} > 0$ 
 $\quad \text{then } tx \{ \text{outputs} = tx.\text{outputs}$ 
 $\quad \cup \{(s, \text{consumed} - \text{produced})\} \text{ else nothing}$ 
 $\text{mkToSpec}' l (\text{MaxFee } f)(tx) = \text{if } tx.\text{fee} \leq f$ 
 $\quad \text{then } tx \text{ else nothing}$ 
 $\text{mkToSpec}' l (\text{AndExps } [a1; a2; \dots; ak])(tx) = \text{mkToSpec}' l \ a1 \ (\dots (\text{mkToSpec}' l \ a2 \ (\text{mkToSpec}' l \ a1 \ tx)))$ 
```

mkToSpec DEFINITION

```
 $\text{mkToSpec} : (\text{LState} \times \mathcal{I}_{\text{post}}) \rightarrow \text{Tx}^?$ 
 $\text{mkToSpec}(l, i) = \text{mkToSpec}' l \ i \ \text{initTx}_{a,mf}$ 
```

Figure 12 Building transactions according to the specific intent

```

txi1 = (outputs={ (s,t) },
           validityInterval=[slot l,(slot l)+j],
           mint=t,
           fee=minfee l
           sigKeys=getSignersVal t)

txi2 = (outputs={(RequireSig k2, x) , (RequireSig k1 ,
           (balance (mkIns l {} x (RequireSig k1)))-x)},
           validityInterval=[ nothing ,nothing ],
           mint={ },
           fee=minfee l
           sigKeys={ k1 })

```

 **Figure 13** Abstract transaction examples